



# Kubernetes

Tecnologias de Virtualização e  
Centros de Dados  
Mestrado em Engenharia Informática

Alexandre Fonte  
[alexandre.fonte@ubi.pt](mailto:alexandre.fonte@ubi.pt)  
(Professor Convidado)  
Departamento de Informática  
Ano Letivo 2023/2024

# Sumário

- Containers
- Kubernetes
- Kubernetes Architecture
- Pods, Services, Replicasets, and Deployments
- Kubernetes Cluster
  - Minikube
  - AKS and others

Versão: 29 abril de 2022,  
Pequenas revisões/atualizações 9 abril de 2023 e 6 abril de 2024.

# Containers

- A tecnologia de containers é cada vez mais popular para o desenvolvimento de sistemas e aplicações
  - Não é necessário configurar e gerir hardware
  - Não é necessário configurar e gerir sistemas operativos
  - Permite poupar tempo
  - Permite poupar recursos computacionais
  - Permite poupar custos
- Os sistemas de gestão de containers normais focam-se em containers individuais. Assim, em grandes projetos, alguns aspetos tornam-se difíceis de gerir:
  - Configuração e gestão de load balancing
  - Conectividade Internet
  - Orquestração do processo de lançamento de containers

# Plataformas de gestão de containers

- De modo a facilitar a gestão de containers em grandes projetos, foram criadas plataformas de gestão de containers
- A mais popular chama-se kubernetes ou K8s
  - Por vezes é abreviada para **K8s**



# kubernetes

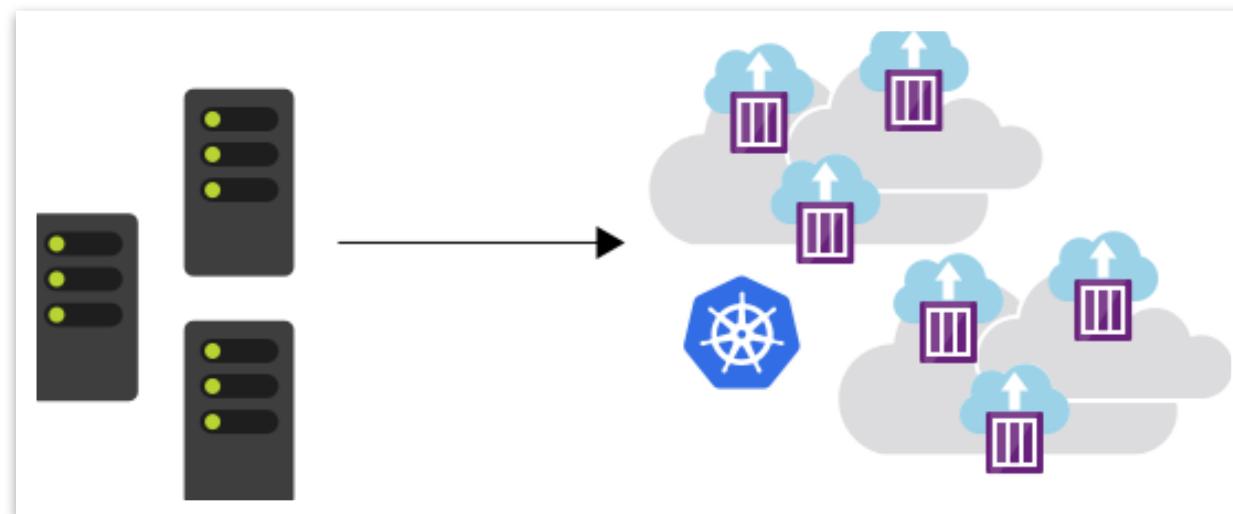
# Kubernetes

- "Kubernetes é um sistema de código aberto para a automatização da implantação, dimensionamento e gestão de aplicações *containerizadas*".
- Kubernetes provém da palavra grega κυβερνήτης, que significa timoneiro ou piloto de navio. Com esta analogia em mente, podemos pensar em Kubernetes como o piloto de um navio de contentores.
- Kubernetes é também referido como k8s (pronuncia-se Kate's), uma vez que existem 8 caracteres entre k e s.
- Kubernetes é altamente inspirado pelo sistema Google Borg, um orquestrador de contentores e carga de trabalho para as suas operações globais.
- É um projecto de código aberto escrito na língua Go e licenciado sob a Licença Apache, Versão 2.0.
- Kubernetes foi iniciado pela Google e, com o seu lançamento v1.0 em Julho de 2015, doou-o à Cloud Native Computing Foundation (CNCF), uma das maiores subfundações da Linux Foundation.
- As novas versões Kubernetes são lançadas em ciclos de 4 meses. A versão estável actual é 1.23 (a partir de Abril de 2022).



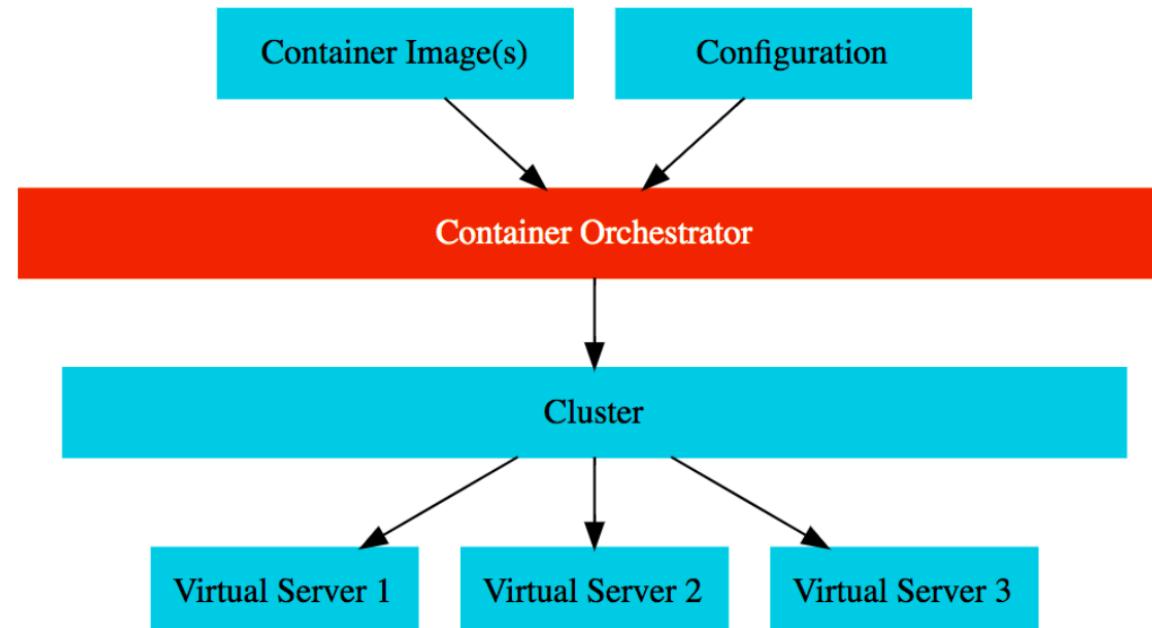
# Kubernetes

- Kubernetes é uma plataforma portável, extensível e de código aberto para gerir e orquestrar aplicações baseadas em containers
- Kubernetes cria uma camada de abstração sobre tarefas complexas de gestão de containers



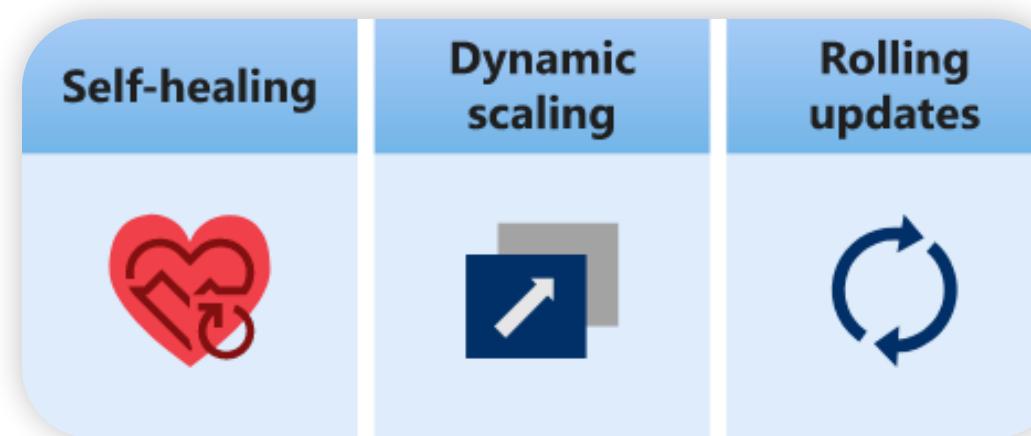
# Kubernetes - Orquestração de Containers

- **Requisito:** Preciso de 10 instâncias do Contentor do Microserviço A, 15 instâncias do Contentor do Microserviço B e ...
- Também preciso de:
  - Auto Scaling** - Escalar containers com base na procura
  - Load Balancing** - Distribuir a carga entre múltiplas instâncias do mesmo microserviço
  - Self Healing** - Verificação da saúde e substituição das instâncias em falha
  - Zero Downtime Deployments** - Atualizações com as novas versões sem downtime



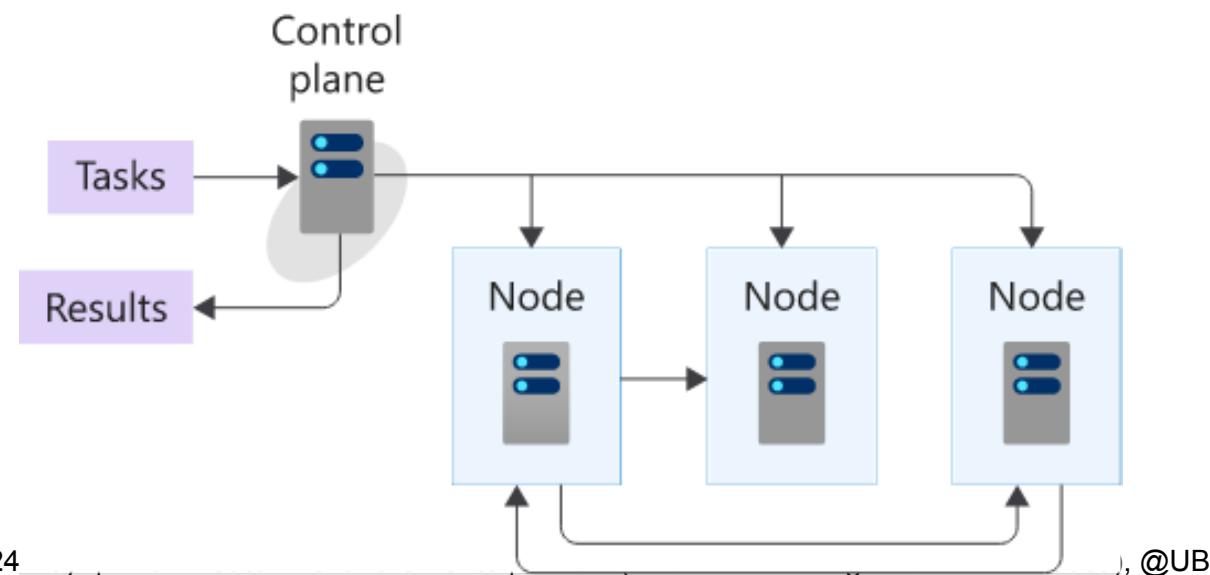
# Vantagens de usar Kubernetes

- Faz a cura automática (restart ou replace) de containers
- Faz escalonamento automático em função da procura
- Automatiza as atualizações de containers
- Faz a gestão de storage
- Faz a gestão de tráfego de rede



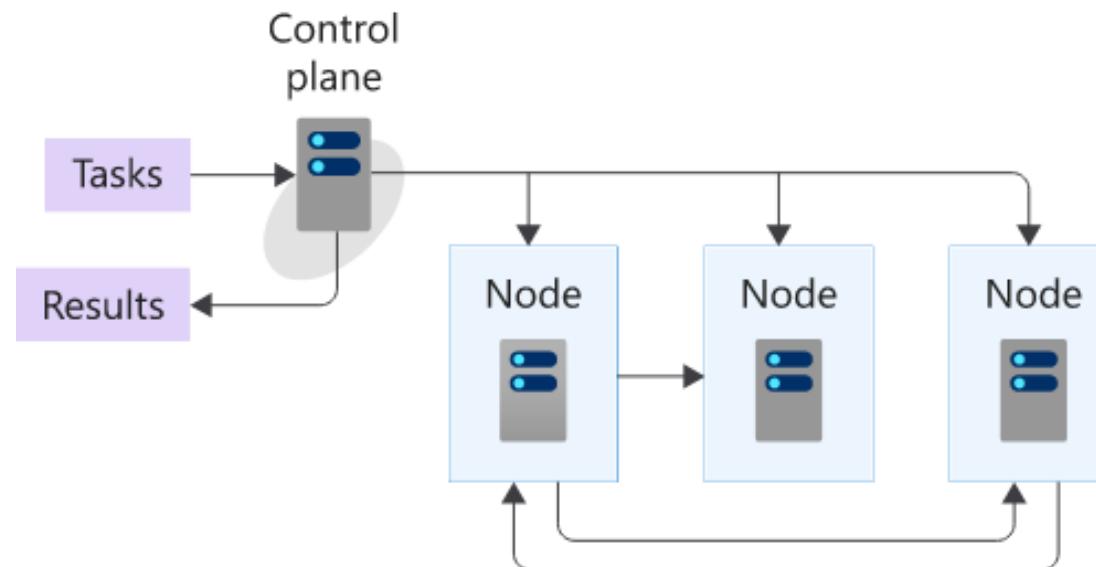
# Cluster

- Um cluster é um conjunto de computadores que trabalham em conjunto e que podem ser vistos como um sistema único
- Cada computador de um cluster chama-se um **nó (node)**
- Os **nós** são tipicamente configurados para executar o mesmo tipo de tarefas, por exemplo alojar websites, correr containers (novo foco), etc



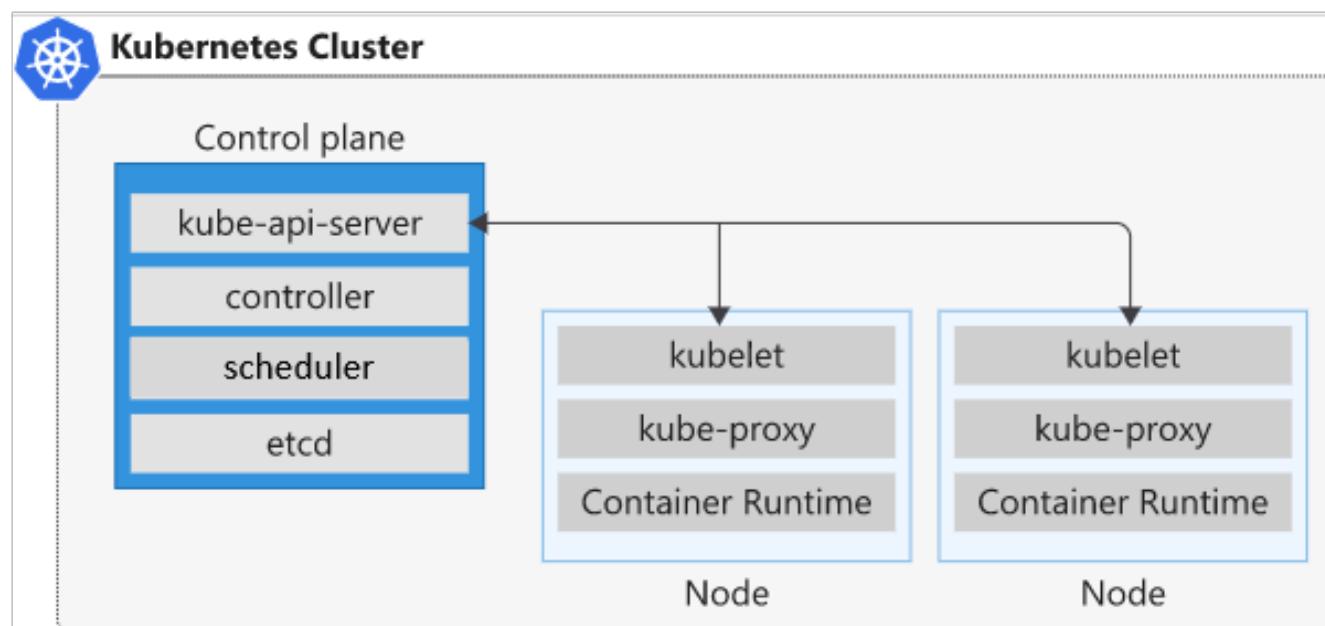
# Cluster

- Um cluster usa software centralizado para agendar e controlar as tarefas que são atribuídas aos **nós**
- Este software centralizado é executado em computadores chamados **control planes**



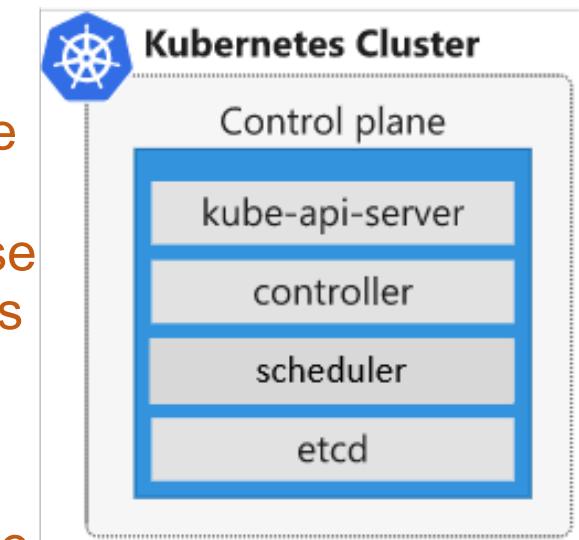
# Arquitetura Kubernetes

- Um cluster Kubernetes tem pelo menos um *control plane* e um nó (que executa as aplicações *containerizadas*)
- O plano de controlo gere os nós e os pods no cluster
- Estes componentes podem correr em máquinas físicas, VM, ou instâncias na Cloud
- O OS predefinido em Kubernetes é o Linux



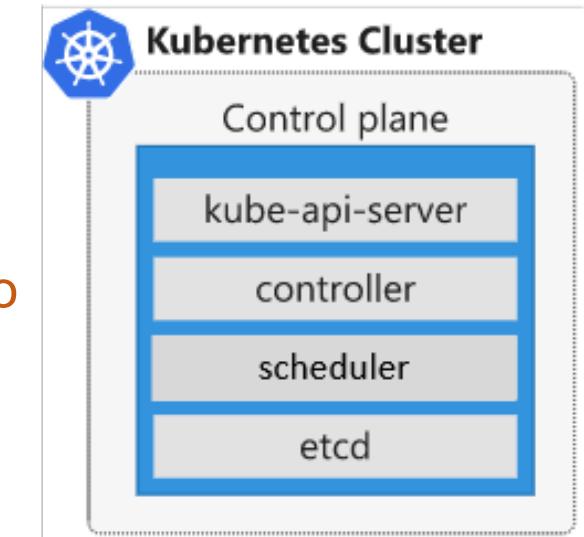
# Componentes do control plane

- **API server:** é o front end de comunicação entre todos os componentes do cluster, através de uma API RESTful
- **Scheduler:** componente responsável pela atribuição de tarefas (Pods) aos nós. Para cada Pod, procura descobrir numa operação a 2-passos o melhor nó para executar esse Pod, de entre os nós admissíveis (i.e, 1.<sup>º</sup>-**Filtering**: filtra os nós com recursos suficientes; 2.<sup>º</sup> **Scoring**: ordena os nós admissíveis com base num score)
- **Backing store (etcd):** sistema de armazenamento de dados principal do cluster (configuração e estado atual do cluster). É um armazenamento do tipo chave/valor.
  - Tudo o que podemos consultar com comando `kubectl get` é armazenado no Etcd.
  - O comando `kubectl create` cria uma nova entrada no etcd



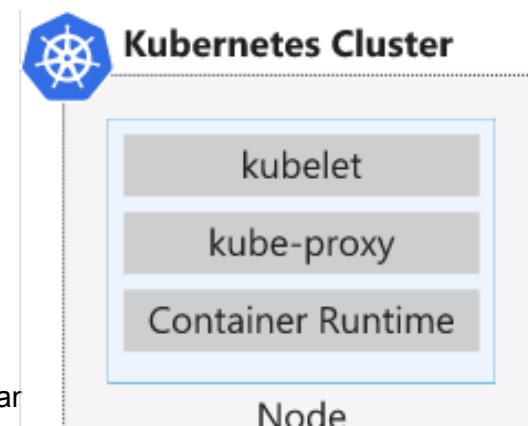
# Componentes do control plane

- **Controller manager:** este componente lança controllers e vigia o estado dos controllers do cluster
- **Controller process:** um controlador é um processo que vigia continuamente o estado dos objetos no cluster, procurando que estes e cluster fiquem no estado desejado (spec field). Existem vários tipos, alguns são:
  - Node controller: responsável por notificar e dar uma resposta quando um nó fica em baixo
  - Job controller: cria Pods para correr as tarefas
- **Cloud controller manager (opcional):** integra kubernetes com as eventuais tecnologias cloud no cluster, que podem ser load balancers, filas ou storage



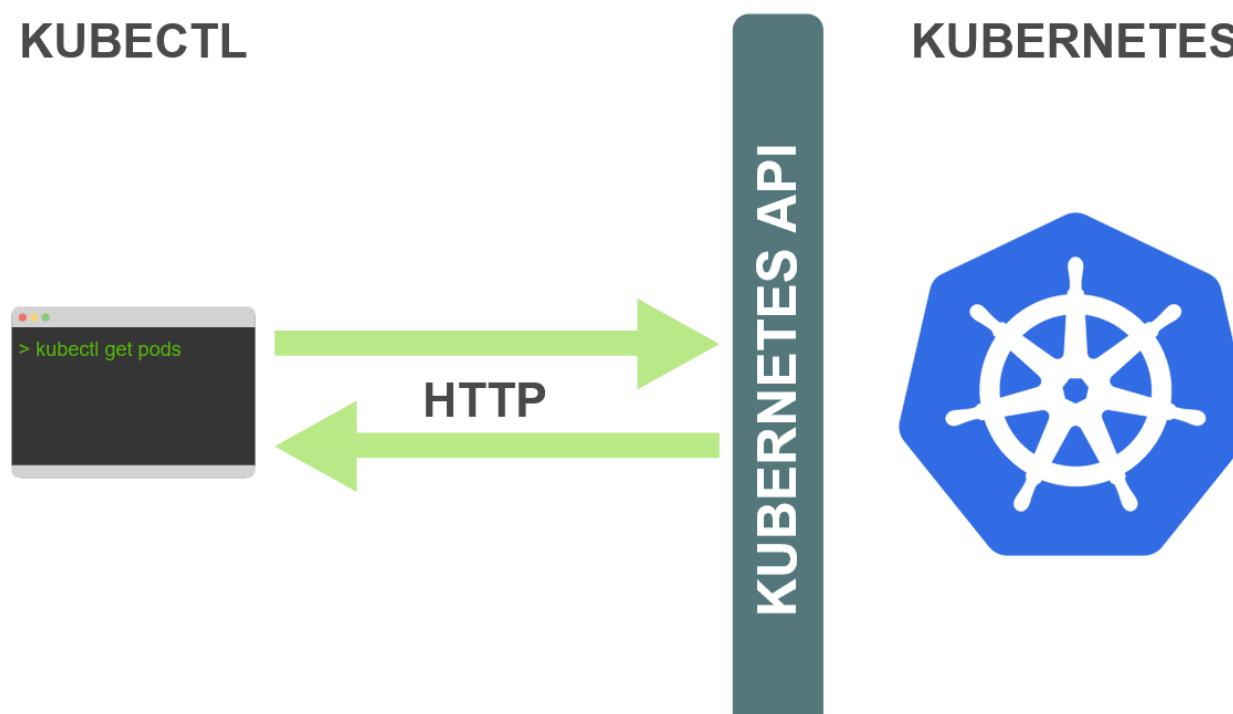
# Componentes do nó

- **Kubelet:** é o agente que corre em cada nó. Este regista cada nó com o API server (pelo hostname; flag; ou outro identificador). Recebe pedidos de tarefas do control plane (através da API) e verifica se as tarefas (os Pods) estão a correr e saudáveis. Mais especificamente, recebe as especificações do Pod (PodSpec) para vigiar os containers descritos no PodSpec.
- **Kube-proxy:** responsável pela configuração de rede do nó, permitindo a comunicação com os Pods através de sessões de rede internas ou externas ao cluster. Também implementa regras para encaminhamento e load balancing. Em súmula implementa parte do conceito de Service.
- **Container runtime:** é o componente responsável por correr os containers. As suas funções incluem obter, iniciar e parar as imagens. Kubernetes suporta vários engines, incluindo Docker, rkt, CRI-O, containerd e frakti.



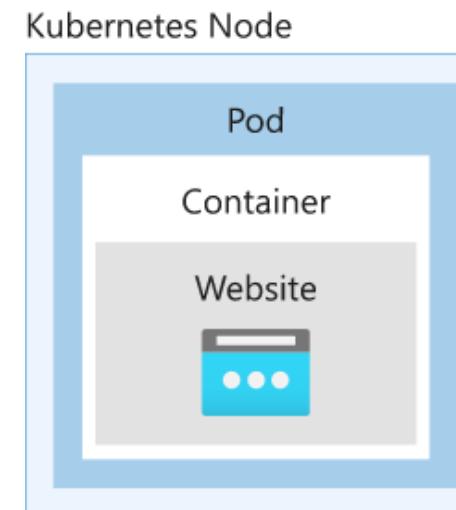
# Usando CLI para interagir com kubernetes

- A interação com um cluster kubernetes pode ser feita através de uma ferramenta de linha de comando chamada **kubectl**
- **kubectl** permite fazer o *deploy* de aplicações, inspecionar e gerir os recursos do cluster e visualizar *logs*.



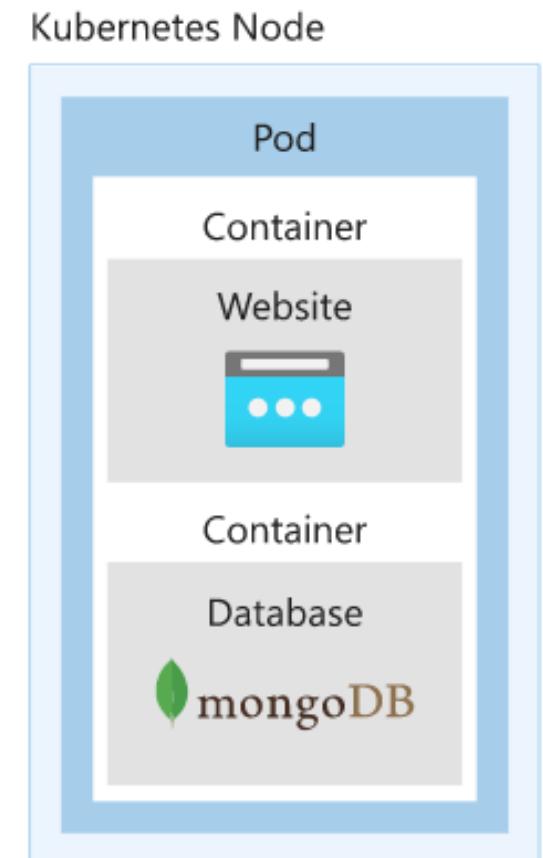
# Kubernetes pod

- Um cluster kubernetes **não consegue correr diretamente containers**
- É preciso primeiro empacotar/encapsular os containers dentro de um objecto chamado **pod**
- O **pod** representa uma instância de uma aplicação a correr num cluster kubernetes
- O **pod** é o objeto mais pequeno que se pode criar em kubernetes

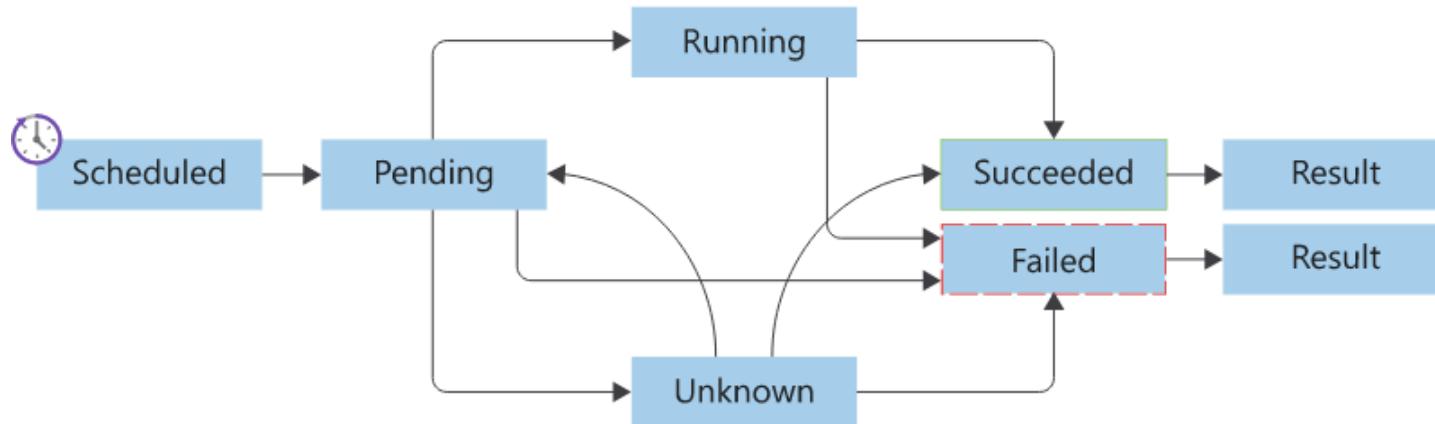


# Kubernetes pod

- Um pod pode conter vários containers
- O pod contém também informação sobre:
  - Storage partilhada
  - Configuração de rede
  - Instruções para correr os containers
- Contudo, se for necessário escalar uma aplicação não é possível criar um novo container no pod. É necessário criar um novo pod no mesmo nó ou outro nó do cluster.



# Pod lifecycle



Pending	After the pod run is scheduled, the container runtime downloads container images, and starts all containers for the pod.
Running	The pod transitions to a running state after all of the resources within the pod are ready.
Succeeded	The pod transitions to a succeeded state after the pod completes its intended task, and runs successfully.
Failed	Pods can fail for various reasons. A container in the pod may have failed, leading to the termination of all other containers. Or, maybe an image wasn't found during preparation of the pod containers. In these types of cases, the pod can go to a failed state. Pods can transition to a failed state from either a pending state or a running state. A specific failure can also place a pod back in the pending state.
Unknown	If the state of the pod can't be determined, the pod is an unknown state.

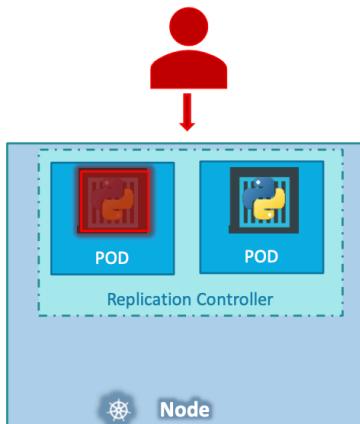
# Pod deployment options com CLI (kubectl)

- **POD templates:** define a configuração do pod através de um ficheiro de configuração. O template tem informação dos nomes das imagens dos containers, o nome do registry a usar para descarregar as imagens, os ports a usar e outros pormenores de runtime.
- **Replication controllers:** usa pod templates e define um determinado número de pods que têm que correr. O replication controller assegura que esse número de pods está sempre a correr, substituindo dinamicamente os que eventualmente falhem
- **Replica sets:** semelhante aos replication controllers (é uma versão moderna), mas adicionalmente consegue identificar pods através de uma etiqueta chamada selector value
- **Deployments:** cria um objeto de gestão acima de um replica set, útil para gerir os updates dos pods num cluster
  - **rolling updates:** lança pods de uma nova versão sem desativar antes os da versão anterior, o que permite zero downtime
  - **rollback:** reverte um update em caso de necessidade, mais uma vez com zero downtime

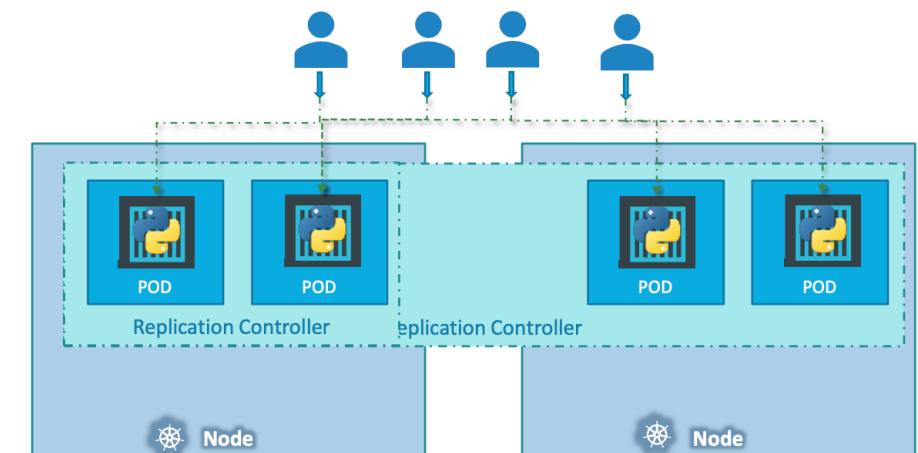
# Replication Controller

- **Replication controllers:** usa pod templates e define um determinado número de pods que têm que correr. O replication controller assegura que esse número de pods está sempre a correr, substituindo dinamicamente os que eventualmente falhem
- 

High Availability (HA)



Load-Balancing and Scaling

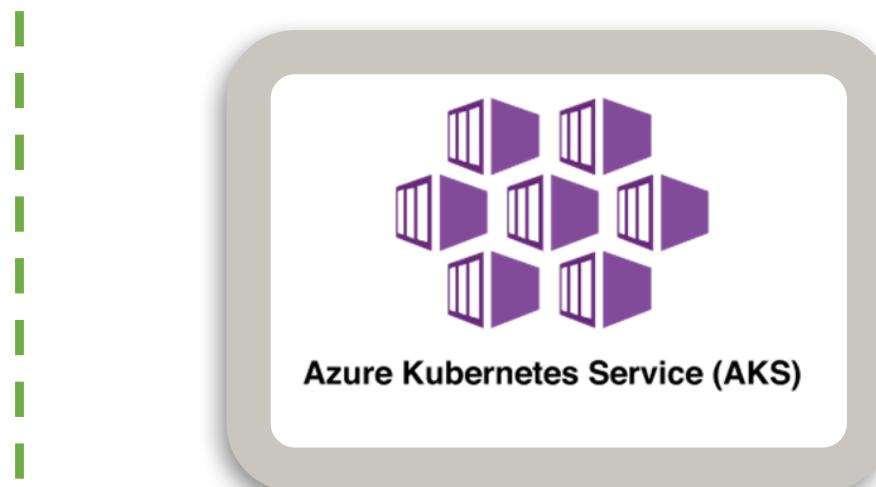


Replicação para assegurar tolerância a falhas.  
Mesmo existindo um único Pod em execução, se este falhar o replication controller recupera-o.

# Como implementar kubernetes

Numa infraestrutura própria    ou    Através de uma plataforma cloud

- É necessário criar o cluster
- Implementar um ambiente de runtime (Docker ou outro)
- Arranjar soluções para load balancing, logging, monitoring
- É necessário gerir as redes virtuais, OS, memória, storage, etc



- O serviço AKS providencia um ambiente kubernetes integrado no Azure, o que reduz a complexidade da sua implementação

# Máquina Local usando o Minikube cluster

- minikube is local Kubernetes, focusing on making it easy to learn and develop for Kubernetes. (<https://minikube.sigs.k8s.io/docs/start/>)
- All you need is Docker (or similarly compatible) container or a Virtual Machine environment, and Kubernetes is a single command away:  
`minikube start`
- Download and install

## Mac OS Systems (Intel x86-64)

```
curl -LO https://storage.googleapis.com/minikube/releases/latest/minikube-darwin-amd64
sudo install minikube-darwin-amd64 /usr/local/bin/minikube
```

## Linux Systems (Intel x86-64)

```
curl -LO https://storage.googleapis.com/minikube/releases/latest/minikube-darwin-amd64
sudo install minikube-darwin-amd64 /usr/local/bin/minikube
```

# Minikube cluster (<https://minikube.sigs.k8s.io/docs/start/>)

- Download and install (cont.)

## Windows Systems via Powershell

```
New-Item -Path 'c:\' -Name 'minikube' -ItemType Directory -Force  
Invoke-WebRequest -OutFile 'c:\minikube\minikube.exe' -Uri 'https://github.com/  
kubernetes/minikube/releases/latest/download/minikube-windows-amd64.exe'  
-UseBasicParsing
```

```
$oldPath = [Environment]::GetEnvironmentVariable('Path',  
[EnvironmentVariableTarget]::Machine)  
if ($oldPath.Split(';') -inotcontains 'C:\minikube'){`  
    [Environment]::SetEnvironmentVariable('Path', $($'{0};C:\minikube' -f $oldPath),  
    [EnvironmentVariableTarget]::Machine)`  
}
```

# Interact with Minikube Cluster

- minikube start
- minikube stop
- minikube pause
  - Pause Kubernetes without impacting deployed applications
- minikube unpause
- minikube dashboard
- minikube dashboard —url
  - Getting just url of the dashboard
- minikube config set memory 16384
  - Increase the default memory limit (requires a restart)

# 1. Deploy an Application into a Cluster

- Deploy and expose an App on Port 8080 using service type: NodeType

-**kubectl create deployment hello-world-rest-api --image=adfente/hello-server**

•**Note:** Creates a new POD with a single container, the name of the deployment in this example is hello-world-rest-api

-**kubectl expose deployment hello-world-rest-api --port=8080 --type=NodePort --name=hello-word-rest-api-service**

•Exposes the service object for an application externally so that clients outside of the Kubernetes cluster can interface with the application

- **How to access this service in Minikube**

-The easiest way to access this service is to let minikube launch a web browser for you (may not work when using minikube on Docker)

•**minikube service hello-world-rest-api**

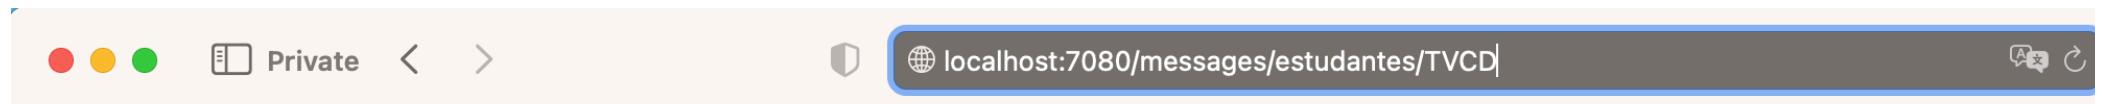
-Alternatively, use kubectl to forward the port:

- **kubectl port-forward <pod name> 7080:8080**
- Tada! Your application is now available at **http://localhost:7080/**
- How to Get Pod name:
  - kubectl get pods**
- Go to the Browser and Enter the following URI

**http://localhost:7080/messages/estudantes/TVCD**

# 1. Deploy an Application into a Cluster (Cont.)

```
[afs-MacBook-Air:~ alexandrefonte$ kubectl get pods
  NAME                  READY   STATUS    RESTARTS   AGE
  hello-world-rest-api-84bf8c5ddf-sftnj   1/1     Running   0          34m
[afs-MacBook-Air:~ alexandrefonte$ kubectl port-forward hello-world-rest-api-84bf8c5ddf-sftnj 7080:8080
Forwarding from 127.0.0.1:7080 -> 8080
Forwarding from [::1]:7080 -> 8080
Handling connection for 7080
```



**Good moorning Dear estudantes from our TVCD Docker Container!**

**It is 13:13, now.**

# Deploy an Application into a Cluster (Cont.)

## - Load Balancing

- Deploy, and Expose on Port 8080 using the LoadBalancer service type

**-kubectl create deployment hello-world-rest-api --image=adfonte/hello-server**

**-kubectl expose deployment hello-world-rest-api --port=8080 --type=LoadBalancer --name=hello-world-rest-service-lb**

**-minikube tunnel**

- connect to LB services (introduzir numa janela à parte) - just in the case of Minikube

**-kubectl get services**

- To check namely the allocated external IP

- Scale a pod deployment to 3 replicas

**-kubectl scale deployment hello-world-rest-api --replicas=3**

# Deploy an Application into a Cluster (Cont.)

## - Load Balancing

- Deploy, and Expose on Port 8080 using the LoadBalancer service type

```
afs-MacBook-Air:~ alexandrefonte$ kubectl expose deployment hello-world-rest-api --port=8080 --type=LoadBalancer --name=hello-world-rest-service-lb
service/hello-world-rest-service-lb exposed
```

```
afs-MacBook-Air:~ alexandrefonte$ kubectl get services
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
hello-world-rest-service-lb	LoadBalancer	10.99.181.183	<pending>	8080:31558/TCP	6s
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	12d
afs-MacBook-Air:~ alexandrefonte\$ kubectl get services					
NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
hello-world-rest-service-lb	LoadBalancer	10.99.181.183	127.0.0.1	8080:31558/TCP	63s
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	12d

In case of  
Minikube running as a VM,  
appears a different IP.  
In a “real” cluster,  
you should get a public  
external IP.



**Good moorning Dear estudantes from our Tecnologias Virtualizacao Docker Container!**

**It is 13:22, now.**

# Deploy an Application into a Cluster (Cont.)

## - Load Balancing and Auto Scaling

- Deploy, and Expose on Port 8080 using the LoadBalancer service type and AutoScaling
  - kubectl create deployment hello-world-rest-api --image=adfonte/hello-server**
  - kubectl expose deployment hello-world-rest-api --port=8080 --type=LoadBalancer --name=hello-world-rest-service-lb**
- AutoScale a pod deployment to 10 replicas
  - kubectl autoscale deployment hello-world-rest-api --max=10 --cpu-percent=70**
  - kubectl autoscale deployment hello-world-rest-api --min=3 --max=10 --cpu-percent=70**

# Deploy an Application into a Cluster (Cont.)

## - Load Balancing and Auto Scaling

- A **HorizontalPodAutoscaler** (HPA) automatically updates a workload resource (such as a Deployment or StatefulSet), with the aim of automatically scaling the workload to match demand.
- Horizontal scaling means that the response to increased load is to deploy more Pods. This is different from vertical scaling, which for Kubernetes would mean assigning more resources (for example: memory or CPU) to the Pods (...)
- If the load decreases, and the number of Pods is above the configured minimum, the HorizontalPodAutoscaler instructs the workload resource to scale back down.

horizontal pod autoscaler status							
NAME	REFERENCE	TARGETS	MINPODS	MAXPODS	REPLICAS	AGE	
hello-world-rest-api	Deployment/hello-world-rest-api	<unknown>/70%	3	10	3	11m	

Precisamos de ativar um metrics-server!  
Precisamos de alguma carga!

# Deploy an Application into a Cluster (Cont.)

## - Load Balancing and Auto Scaling

- Ativação de um metrics-server no minikube k8s

### **-minikube addons enable metrics-server**

```
[afs-MacBook-Air:~ alexandrefonte$ minikube addons enable metrics-server
  ■ Using image k8s.gcr.io/metrics-server/metrics-server:v0.4.2
  ★ The 'metrics-server' addon is enabled
```

ADDON NAME	PROFILE	STATUS	MAINTAINER
ambassador	minikube	disabled	third-party (ambassador)
auto-pause	minikube	disabled	google
csi-hostpath-driver	minikube	disabled	kubernetes
dashboard	minikube	enabled <input checked="" type="checkbox"/>	kubernetes
default-storageclass	minikube	enabled <input checked="" type="checkbox"/>	kubernetes
efk	minikube	disabled	third-party (elastic)
freshpod	minikube	disabled	google
gcp-auth	minikube	disabled	google
gvisor	minikube	disabled	google
helm-tiller	minikube	disabled	third-party (helm)
ingress	minikube	disabled	unknown (third-party)
ingress-dns	minikube	disabled	google
istio	minikube	disabled	third-party (istio)
istio-provisioner	minikube	disabled	third-party (istio)
kong	minikube	disabled	third-party (Kong HQ)
kubevirt	minikube	disabled	third-party (kubevirt)
logviewer	minikube	disabled	unknown (third-party)
metallb	minikube	disabled	third-party (metallb)
metrics-server	minikube	disabled	kubernetes
nvidia-driver-installer	minikube	disabled	google
nvidia-gpu-device-plugin	minikube	disabled	third-party (nvidia)
olm	minikube	disabled	third-party (operator)
			framework)
pod-security-policy	minikube	disabled	unknown (third-party)
portainer	minikube	disabled	portainer.io
registry	minikube	disabled	google
registry-aliases	minikube	disabled	unknown (third-party)
registry-creds	minikube	disabled	third-party (upmc enterprises)
storage-provisioner	minikube	enabled <input checked="" type="checkbox"/>	google
storage-provisioner-gluster	minikube	disabled	unknown (third-party)
volumesnapshots	minikube	disabled	kubernetes

ADDON NAME	PROFILE	STATUS	MAINTAINER
ambassador	minikube	disabled	third-party (ambassador)
auto-pause	minikube	disabled	google
csi-hostpath-driver	minikube	disabled	kubernetes
dashboard	minikube	enabled <input checked="" type="checkbox"/>	kubernetes
default-storageclass	minikube	enabled <input checked="" type="checkbox"/>	kubernetes
efk	minikube	disabled	third-party (elastic)
freshpod	minikube	disabled	google
gcp-auth	minikube	disabled	google
gvisor	minikube	disabled	google
helm-tiller	minikube	disabled	third-party (helm)
ingress	minikube	disabled	unknown (third-party)
ingress-dns	minikube	disabled	google
istio	minikube	disabled	third-party (istio)
istio-provisioner	minikube	disabled	third-party (istio)
kong	minikube	disabled	third-party (Kong HQ)
kubevirt	minikube	disabled	third-party (kubevirt)
logviewer	minikube	disabled	unknown (third-party)
metallb	minikube	disabled	third-party (metallb)
metrics-server	minikube	enabled <input checked="" type="checkbox"/>	kubernetes
nvidia-driver-installer	minikube	disabled	google
nvidia-gpu-device-plugin	minikube	disabled	third-party (nvidia)
olm	minikube	disabled	third-party (operator)
			framework)
pod-security-policy	minikube	disabled	unknown (third-party)
portainer	minikube	disabled	portainer.io
registry	minikube	disabled	google
registry-aliases	minikube	disabled	unknown (third-party)
registry-creds	minikube	disabled	third-party (upmc enterprises)
storage-provisioner	minikube	enabled <input checked="" type="checkbox"/>	google
storage-provisioner-gluster	minikube	disabled	unknown (third-party)
volumesnapshots	minikube	disabled	kubernetes

# Deploy an Application into a Cluster (Cont.)

## - Load Balancing and Auto Scaling

- Verificação se o metrics-server está em execução no Cluster

**-kubectl get pods --all-namespaces | grep metrics-server**

```
● ● ● alexandrefonte — -bash — 107x24
[af5-MacBook-Air-2:~ alexandrefonte$ kubectl get pods --all-namespaces | grep metrics-server
kube-system           metrics-server-6588d95b98-rcdgk      1/1     Running   0          17m
af5-MacBook-Air-2:~ alexandrefonte$
```

- Consultar as métricas através da metrics API

**-# Get the metrics for all nodes**

**-kubectl get --raw /apis/metrics.k8s.io/v1beta1/nodes**

**-# Get the metrics for all pods**

**-kubectl get --raw /apis/metrics.k8s.io/v1beta1/pods**

- Consultar os pods “Once it is enabled, you can view the metrics of a running pod using the following command”:

**-kubectl top pods**

```
● ● ● alexandrefonte — -bash — 127x24
[af5-MacBook-Air-2:~ alexandrefonte$ kubectl top pod
NAME                  CPU(cores)   MEMORY(bytes)
hello-world-rest-api-6fbb455794-9dpkd  4m        189Mi
hello-world-rest-api-6fbb455794-bg7mk  4m        206Mi
hello-world-rest-api-6fbb455794-lcv5g  4m        219Mi
af5-MacBook-Air-2:~ alexandrefonte$
```

# Deploy an Application into a Cluster (Cont.)

## - Load Balancing and Auto Scaling

- Algoritmo:

```
desiredReplicas = ceil[currentReplicas * ( currentMetricValue / desiredMetricValue )]
```

Exemplo:

$$\text{desiredReplicas} = 1 * 305 / 50 = 6.1$$

(Arredondado para cima)

NAME	REFERENCE	TARGET	MINPODS	MAXPODS	REPLICAS	AGE
php-apache	Deployment/php-apache/scale	305% / 50%	1	10	1	3m

NAME	REFERENCE	TARGET	MINPODS	MAXPODS	REPLICAS	AGE
php-apache	Deployment/php-apache/scale	305% / 50%	1	10	7	3m

# Deploy an Application into a Cluster (Cont.)

## - Load Balancing and Auto Scaling

```

afs-MacBook-Air:~ alexandrefonte$ kubectl autoscale deployment hello-world-rest-api --max=10 --cpu-percent=70
[horizontalpodautoscaler.autoscaling/hello-world-rest-api autoscaled
afs-MacBook-Air:~ alexandrefonte$ kubectl autoscale deployment hello-world-rest-api --min=3 --max=10 --cpu-percent=70
Error from server (AlreadyExists): horizontalpodautoscalers.autoscaling "hello-world-rest-api" already exists
afs-MacBook-Air:~ alexandrefonte$ kubectl get all
NAME           READY   STATUS    RESTARTS   AGE
pod/hello-world-rest-api-84bf8c5ddf-hr7rz   1/1     Running   0          14m

NAME             TYPE        CLUSTER-IP      EXTERNAL-IP      PORT(S)        AGE
service/hello-world-rest-service-lb   LoadBalancer   10.99.181.183  127.0.0.1       8080:31558/TCP  13m
service/kubernetes   ClusterIP      10.96.0.1       <none>          443/TCP       12d

NAME           READY   UP-TO-DATE   AVAILABLE   AGE
deployment.apps/hello-world-rest-api   1/1     1           1          14m

NAME           DESIRED   CURRENT   READY   AGE
replicaset.apps/hello-world-rest-api-84bf8c5ddf   1         1         1        14m
[

NAME           REFERENCE          TARGETS   MINPODS   MAXPODS   REPLICAS   AGE
[horizontalpodautoscaler.autoscaling/hello-world-rest-api   Deployment/hello-world-rest-api  <unknown>/70%  1          10          1        47s
afs-MacBook-Air:~ alexandrefonte$ kubectl delete horizontalpodautoscaler.autoscaling/hello-world-rest-api
[horizontalpodautoscaler.autoscaling "hello-world-rest-api" deleted
afs-MacBook-Air:~ alexandrefonte$ kubectl autoscale deployment hello-world-rest-api --min=3 --max=10 --cpu-percent=70
horizontalpodautoscaler.autoscaling/hello-world-rest-api autoscaled
afs-MacBook-Air:~ alexandrefonte$ kubectl get all
NAME           READY   STATUS    RESTARTS   AGE
pod/hello-world-rest-api-84bf8c5ddf-422lg   1/1     Running   0          28s
pod/hello-world-rest-api-84bf8c5ddf-6zphf   1/1     Running   0          28s
pod/hello-world-rest-api-84bf8c5ddf-hr7rz   1/1     Running   0          16m

NAME             TYPE        CLUSTER-IP      EXTERNAL-IP      PORT(S)        AGE
service/hello-world-rest-service-lb   LoadBalancer   10.99.181.183  127.0.0.1       8080:31558/TCP  14m
service/kubernetes   ClusterIP      10.96.0.1       <none>          443/TCP       12d

NAME           READY   UP-TO-DATE   AVAILABLE   AGE
deployment.apps/hello-world-rest-api   3/3     3           3          16m

NAME           DESIRED   CURRENT   READY   AGE
replicaset.apps/hello-world-rest-api-84bf8c5ddf   3         3         3        16m
[

NAME           REFERENCE          TARGETS   MINPODS   MAXPODS   REPLICAS   AGE
horizontalpodautoscaler.autoscaling/hello-world-rest-api   Deployment/hello-world-rest-api  <unknown>/70%  3          10          3        34s

```

# Get Pods, Services, Replicasets and Deployments

- kubectl get services *hello-world-rest-api*
- kubectl get pods
  - To check that an application deployment create one or more pods
- Kubectl get describe pods
  - To get more detailed info about pods
- kubectl get replicasets
- Kubectl get nodes
  - To get the list and status about nodes in a cluster
- Kubectl get describe nodes <*name of the node*>
  - To get more detailed information about a node

Complete reference commands:

<https://kubernetes.io/docs/reference/kubectl/>

# Deleting Pods, Services or Deployments

- **kubectl delete pod *hello-world-rest-api-pod***
- **kubectl delete service *hello-world-rest-api***
- **kubectl delete deployment *hello-world-rest-api***

Complete reference commands:

<https://kubernetes.io/docs/reference/kubectl/>

# Create a Pod with Yaml

- pod-definition.yaml

```
apiVersion: v1
kind: Pod
metadata:
  name: myapp-pod
  labels:
    app: myapp
    type: front-end
spec:
  containers:
    - name: nginx-container
      image: nginx:1.7.0
      ports:
        - containerPort: 80
    - name: postgres
      image: postgres
      env:
        - name: POSTGRES_PASSWORD
          value: mysecretpassword
```

## Pod Template

**kind:** can be a pod, service, replicaSet, deployment

**metadata:** is a dictionary with the data about the object

**labels:** is a dictionary with several data defined freely

**containers:** is a List/Array of the pod containers

**spec:** is a dictionary with the specification of the pod

```
kubectl create -f pod-definition.yml
Kubectl get pods
```

# Atividade - Create a Pod with Yaml

- Create a Pod named **hello-word-rest-api-pod** based on **adfonte/hello-server** image. Expose Pod on Port 8080.

## -Manually using the **kubectl run** command

```
[afs-MacBook-Air:kubernetes alexandrefonte$ kubectl run hello-world-rest-api-pod --image=adfonte/hello-server
pod/hello-world-rest-api-pod created
[afs-MacBook-Air:kubernetes alexandrefonte$ kubectl get pods
NAME                  READY   STATUS    RESTARTS   AGE
hello-world-rest-api-pod 1/1     Running   0          19s
```

## -With a Yaml file

```
apiVersion: v1
kind: Pod
metadata:
  name: hello-world-rest-api-pod
  labels:
    app: hello-world-rest-api
    type: front-end
spec:
  containers:
    - name: hello-world-rest-api
      image: adfonte/hello-server
      ports:
        - containerPort: 8080
```

```
[afs-MacBook-Air:kubernetes alexandrefonte$ kubectl create -f pod-definition.yaml
pod/hello-world-rest-api-pod created
[afs-MacBook-Air:kubernetes alexandrefonte$ kubectl get pods
NAME                  READY   STATUS    RESTARTS   AGE
hello-world-rest-api-pod 0/1     ContainerCreating   0          3s
```

# Create a Pod with Yaml

- Result:

```
[afs-MacBook-Air:kubernetes alexandrefonte]$ kubectl expose pod hello-world-rest-api-pod --port=8080 --type=LoadBalancer --name=hello-world-service-lb
service/hello-world-service-lb exposed
[afs-MacBook-Air:kubernetes alexandrefonte]$ kubectl get services
NAME           TYPE      CLUSTER-IP   EXTERNAL-IP   PORT(S)          AGE
hello-world-service-lb   LoadBalancer   10.108.105.231  127.0.0.1    8080:32168/TCP   2m1s
kubernetes     ClusterIP   10.96.0.1    <none>        443/TCP         7d5h
```

In case of  
Minikube running as a VM,  
appears a different IP

# Create a Replication Controller with Yaml

- rc-definition.yaml

```
apiVersion: v1
kind: ReplicationController
metadata:
  name: myapp-rc
  labels:
    app: myapp
    type: front-end
spec:
  template:
    metadata:
      name: myapp-pod
      labels:
        app: myapp
        type: front-end
    spec:
      containers:
        - name: nginx-container
          image: nginx:1.7.0
          ports:
            - containerPort: 80
        - name: postgres
          image: postgres
          env:
            - name: POSTGRES_PASSWORD
              value: mysecretpassword
replicas: 3
```

## Pod Template

**spec**: now it includes the specification of pods

**template**: the template section declares pods

**replicas**: declares the number of replicas of pods

```
kubectl create -f rc-definition.yaml
kubectl get replicationcontroller
Kubectl get pods
```

# Atividade - Create a Replication Controller with Yaml

- Create a Replication Controller named **hello-word-rest-api-rc** based on adfonte/hello-server image. Create 3 replicas.
- Expose replication controller on Port 8080 using a LoadBalancer.

```
apiVersion: v1
kind: ReplicationController
metadata:
  name: hello-world-rest-api-rc
  labels:
    app: hello-world-rest-api
    type: front-end
spec:
  template:
    metadata:
      name: hello-world-rest-api-pod
      labels:
        app: hello-world-rest-api
        type: front-end
    spec:
      containers:
        - name: hello-world-rest-api
          image: adfonte/hello-server
          ports:
            - containerPort: 8080
      replicas: 3
```

# Create a Replication Controller with Yaml

- Result:

```
replicationcontroller/hello-world-rest-api-rc created
[af5-MacBook-Air:kubernetes alexandrefonte$ kubectl get pods
NAME           READY   STATUS    RESTARTS   AGE
hello-world-rest-api-rc-g5d5f  1/1     Running   0          16s
hello-world-rest-api-rc-l6p5c  1/1     Running   0          16s
hello-world-rest-api-rc-rhcrp  1/1     Running   0          16s
[af5-MacBook-Air:kubernetes alexandrefonte$ kubectl get all
NAME           READY   STATUS    RESTARTS   AGE
pod/hello-world-rest-api-rc-g5d5f  1/1     Running   0          33s
pod/hello-world-rest-api-rc-l6p5c  1/1     Running   0          33s
pod/hello-world-rest-api-rc-rhcrp  1/1     Running   0          33s
NAME           DESIRED  CURRENT  READY   AGE
replicationcontroller/hello-world-rest-api-rc  3        3       3      33s
NAME           TYPE      CLUSTER-IP   EXTERNAL-IP   PORT(S)   AGE
service/kubernetes  ClusterIP  10.96.0.1   <none>        443/TCP  7d5h
[af5-MacBook-Air:kubernetes alexandrefonte$ kubectl expose replicationcontroller hello-world-rest-api-rc --port=8080 --type=LoadBalancer --name=hello-world-service-lb
service/hello-world-service-lb exposed
[af5-MacBook-Air:kubernetes alexandrefonte$ kubectl get services
NAME           TYPE      CLUSTER-IP   EXTERNAL-IP   PORT(S)   AGE
hello-world-service-lb  LoadBalancer  10.111.135.185  127.0.0.1    8080:31185/TCP  5m43s
kubernetes     ClusterIP  10.96.0.1   <none>        443/TCP  7d5h
```

In case of  
Minikube running as a VM,  
appears a different IP

Now Scale to 6 replicas:

```
[af5-MacBook-Air:kubernetes alexandrefonte$ kubectl scale --replicas=6 -f rc-definition.yaml
replicationcontroller/hello-world-rest-api-rc scaled
[af5-MacBook-Air:kubernetes alexandrefonte$ kubectl get replicationcontroller
NAME           DESIRED  CURRENT  READY   AGE
hello-world-rest-api-rc  6        6       6      11m
```

# Create a Replicaset with Yaml

- replicaset-definition.yaml

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: myapp-replicaset
  labels:
    app: myapp
    tier: front-end
spec:
```

## Pod Template

```
template:
  metadata:
    name: myapp-pod
    labels:
      app: myapp
      tier: front-end
  spec:
    containers:
      - name: nginx-container
        image: nginx:1.7.0
        ports:
          - containerPort: 80
      - name: postgres
        image: postgres
        env:
          - name: POSTGRES_PASSWORD
            value: mysecretpassword
    replicas: 3
    selector:
      matchLabels:
        tier: front-end
```

**apiVersion:** apps/v1 this is the api version that supports ReplicaSet

```
kubectl create -f replicaset-definition.yaml
kubectl get replicaset
Kubectl get pods
Kubectl delete replicaset myapp-replicaset
```

**selector:** enable replicaset to identify the pods to monitor by filtering them by their labels.  
**Example:** Filtering by tier label.

# Actividade - Create a Replicaset with Yaml

- Create a ReplicationSet named **hello-word-rest-api-rs** based on adfonte/hello-server image. Create 3 replicas. Filter pods by tier label front-end
- Expose **replicationset** on Port 8080 using a LoadBalancer.

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: hello-world-rest-api-rs
  labels:
    app: hello-world-rest-api
    tier: front-end
spec:
  replicas: 3
  selector:
    matchLabels:
      tier: front-end
  template:
    metadata:
      name: hello-world-rest-api-pod
      labels:
        app: hello-world-rest-api
        tier: front-end
    spec:
      containers:
        - name: hello-world-rest-api
          image: adfonte/hello-server
          ports:
            - containerPort: 8080
```

# Create a Replicaset with Yaml

- Result:

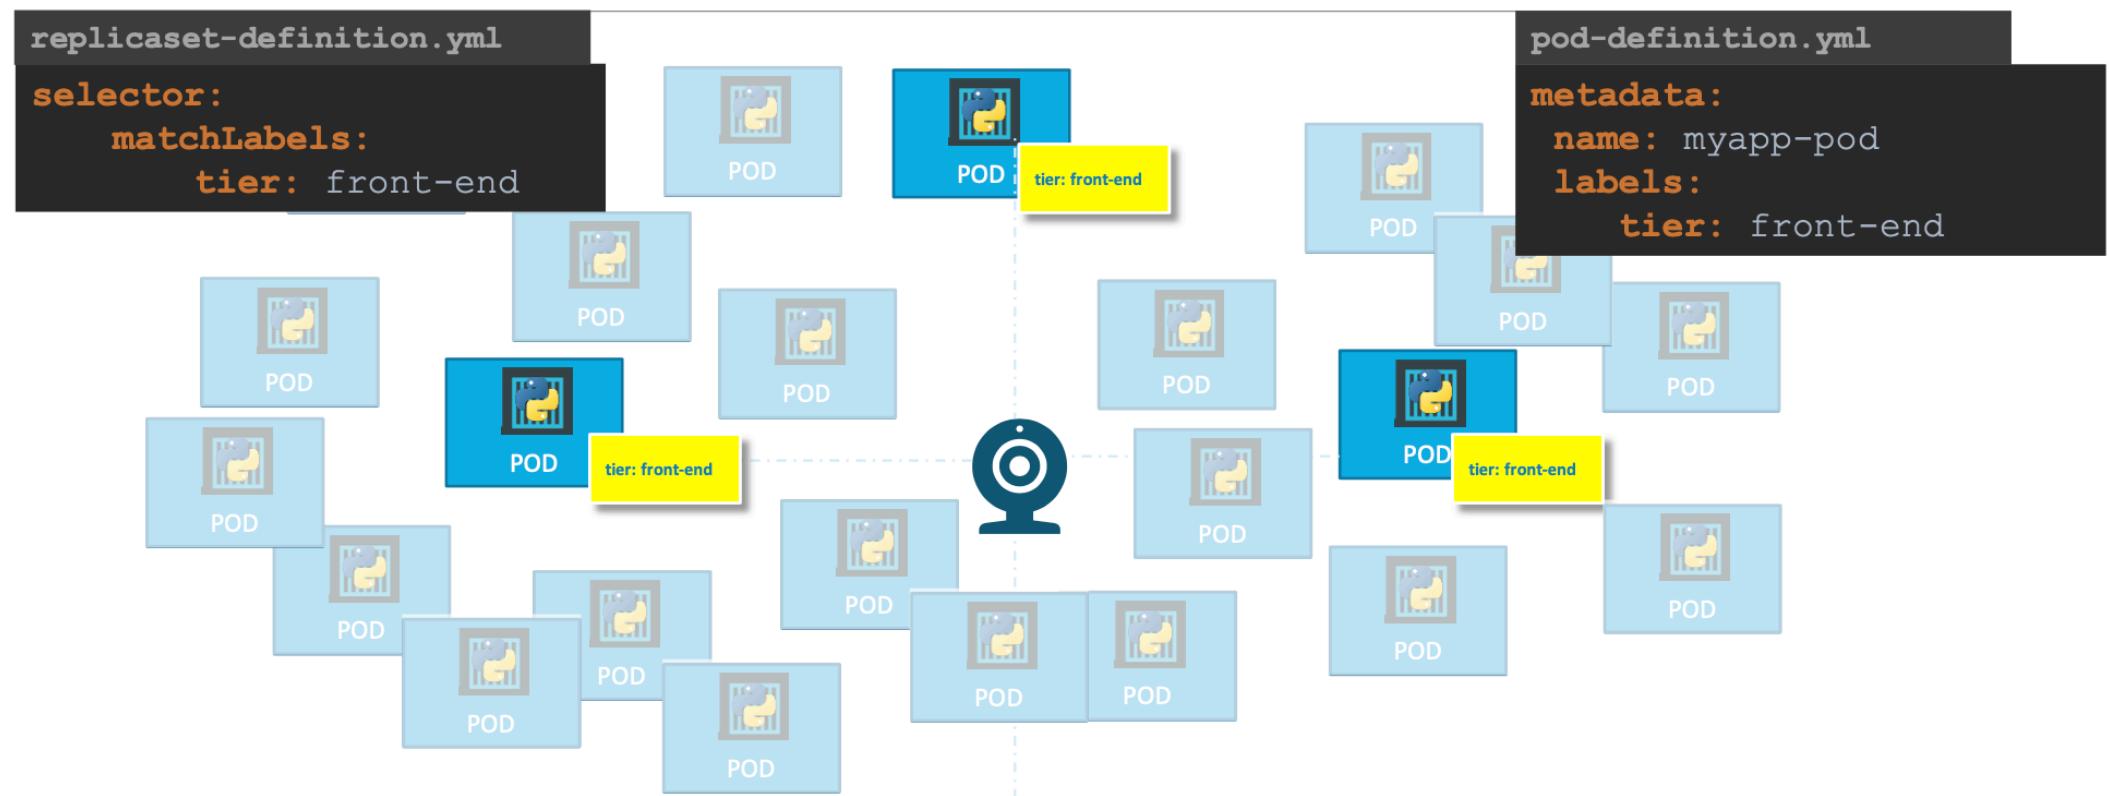
```
[afs-MacBook-Air:kubernetes alexandrefonte$ kubectl create -f replicationset-definition.yaml
replicaset.apps/hello-world-rest-api-rs created
[afs-MacBook-Air:kubernetes alexandrefonte$ kubectl create -f replicationset-definition.yaml
Error from server (AlreadyExists): error when creating "replicationset-definition.yaml": replicasets.apps "hello-world-rest-api-rs" already exists
[afs-MacBook-Air:kubernetes alexandrefonte$ kubectl get all
NAME                           READY   STATUS    RESTARTS   AGE
pod/hello-world-rest-api-rs-5drv  1/1     Running   0          4m13s
pod/hello-world-rest-api-rs-bzqnk  1/1     Running   0          4m13s
pod/hello-world-rest-api-rs-fgzmj  1/1     Running   0          4m13s

NAME            TYPE        CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
service/kubernetes   ClusterIP   10.96.0.1    <none>           443/TCP    7d6h

NAME                DESIRED   CURRENT   READY   AGE
replicaset.apps/hello-world-rest-api-rs  3         3         3       4m13s
```

```
[afs-MacBook-Air:kubernetes alexandrefonte$ kubectl create -f replicationset-definition.yaml
replicaset.apps/hello-world-rest-api-rs created
[afs-MacBook-Air:kubernetes alexandrefonte$ kubectl get replicasets
NAME                DESIRED   CURRENT   READY   AGE
hello-world-rest-api-rs  3         3         3       9s
[afs-MacBook-Air:kubernetes alexandrefonte$ kubectl expose replicaset hello-world-rest-api-rs --port=8080 --type=LoadBalancer --name=hello-world-service-lb
service/hello-world-service-lb exposed
```

# Create a Replicaset with Yaml



# How to Scale

- **Way 1 - By editing the replicaset-definition.yaml**

- Update in the file the section **replicas: 6**

- Kubectl replace -f replicaset-definition.yaml**

- **Way 2 - By using the kubectl scale command**

- Kubectl scale --replicas=6 -f replicaset-definition.yaml**

- Nota: Mas não atualiza o conteúdo do ficheiro.

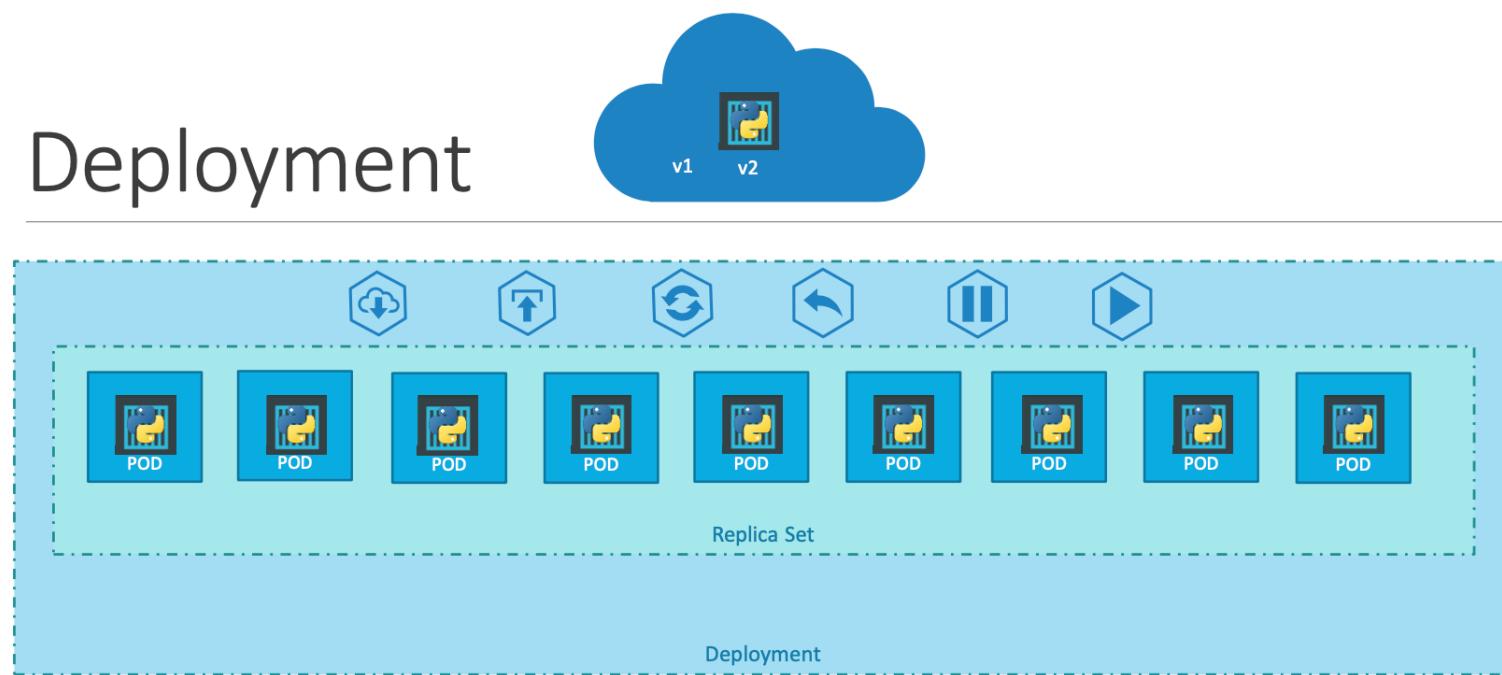
- OR

- Kubectl scale --replicas=6 replicaset.apps/hello-world-rest-api-rs**

- **Way 3 - Automatically based on load (autoscale command)**

# Kubernetes Deployments

- PODs deploy single instances of our application such as the web application. Each container of our is encapsulated in PODs.
- Multiple such PODs are deployed using Replication Controllers or Replica Sets
- Deployment is a kubernetes object that comes higher in the hierarchy.
- The deployment provides us with capabilities to upgrade the underlying instances seamlessly using rolling updates, undo changes (in case of unexpected error during an update), and pause and resume changes to deployments



# Create a Deployment with Yaml

- deployment-definition.yaml

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: myapp-deployment
  labels:
    app: myapp
    tier: front-end
spec:
  template:
    metadata:
      name: mysql:5.7
      labels:
        app: myapp
        tier: front-end
    spec:
      containers:
        - name: nginx-container
          image: nginx:1.7.0
          ports:
            - containerPort: 80
        - name: postgres
          image: postgres
          env:
            - name: POSTGRES_PASSWORD
              value: mysecretpassword
      replicas: 3
      selector:
        matchLabels:
          tier: front-end
```

**kind:** deployment.  
The rest is identically  
to a replicaset

```
kubectl create -f deployment-definition.yaml
kubectl get deployments
kubectl get replicasets
Kubectl get pods
```

# Atividade - Create a Deployment with Yaml

- deployment-definition.yaml
- Create a Deployment named **hello-world-rest-api-deployment** based previous replicaset. Define below replicas section a label with the strategy type: RollingUpdate
- Expose **deployment** on Port 8080 using a LoadBalancer.

```
[afs-MacBook-Air:kubernetes alexandrefonte$ kubectl create -f deployment-definition.yaml
deployment.apps/hello-world-rest-api-deployment created
[afs-MacBook-Air:kubernetes alexandrefonte$ kubectl get all
NAME                                         READY   STATUS            RESTARTS   AGE
pod/hello-world-rest-api-deployment-7d645f687f-9fr2q   0/1    ContainerCreating   0          4s
pod/hello-world-rest-api-deployment-7d645f687f-jb7tz   0/1    ContainerCreating   0          4s
pod/hello-world-rest-api-deployment-7d645f687f-nrsxr   1/1    Running           0          4s

NAME          TYPE        CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
service/kubernetes   ClusterIP   10.96.0.1    <none>           443/TCP     7d6h

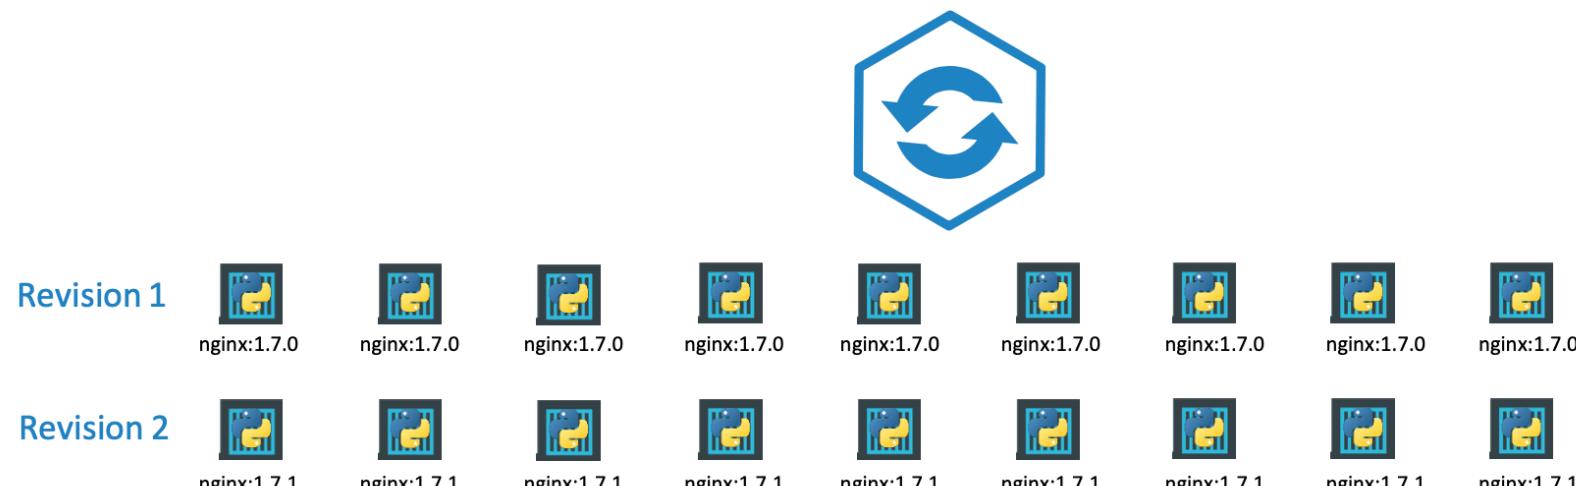
NAME                                         READY   UP-TO-DATE   AVAILABLE   AGE
deployment.apps/hello-world-rest-api-deployment   1/3     3           1           7s

NAME                                         DESIRED   CURRENT   READY   AGE
replicaset.apps/hello-world-rest-api-deployment-7d645f687f   3         3         1       6s
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: hello-world-rest-api-deployment
  labels:
    app: hello-world-rest-api
    tier: front-end
spec:
  replicas: 3
  strategy:
    type: RollingUpdate
  selector:
    matchLabels:
      tier: front-end
  template:
    metadata:
      name: hello-world-rest-api-pod
      labels:
        app: hello-world-rest-api
        tier: front-end
    spec:
      containers:
        - name: hello-world-rest-api
          image: adfonte/hello-server
          ports:
            - containerPort: 8080
```

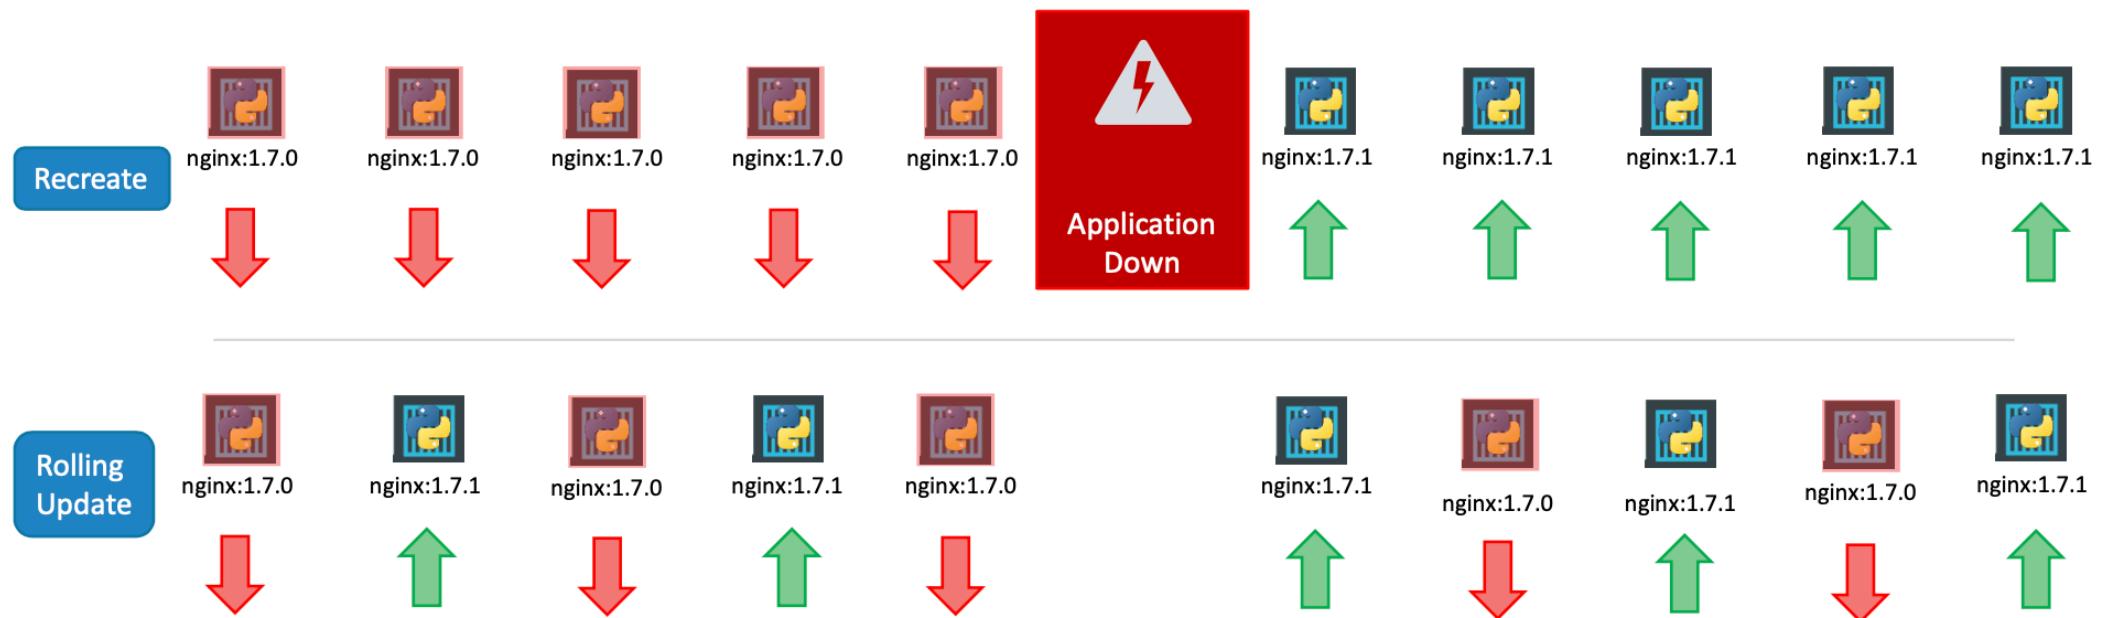
# Rollout and Versioning

- A rollout is the process of gradually deploying or upgrading your application containers.
- When you first create a deployment, it triggers a rollout. A new rollout creates a new Deployment revision. Let's call it revision 1.
- In the future when the application is upgraded – meaning when the container version is updated to a new one – a new rollout is triggered and a new deployment revision is created named Revision 2.
- This helps us keep track of the changes made to our deployment and enables us to rollback to a previous version of deployment if necessary.



# Types of Deployment Strategies

- **Recreate** - First all instances are destroyed, a then are deployed all new instances of the new application version
- **Rolling Update (Default)**- The instances are destroyed and then deployed new instances, one by one.



# How to update a Deployment

- deployment-definition.yml

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: myapp-deployment
  labels:
    app: myapp
    type: front-end
spec:
  template:
    metadata:
      name: mysql:5.7
      labels:
        app: myapp
        type: front-end
    spec:
      containers:
        - name: nginx-container
          image: nginx:latest
          ports:
            - containerPort: 80
        - name: postgres
          image: postgres
          env:
            - name: POSTGRES_PASSWORD
              value: mysecretpassword
      replicas: 3
      selector:
        matchLabels:
          Tier: front-end
```



Update with:

```
kubectl apply -f deployment-definition.yml
```

Or

**kind**: deployment.  
The rest is identically  
to a replicaset

Warning: The deployment-definition  
File must be then updated  
with the new configuration

```
kubectl set image deployment/myapp-deployment nginx=nginx:latest
```

To check status or history:

```
Kubectl describe deployment deployment/myapp-deployment
```

```
Kubectl rollout status deployment deployment/myapp-deployment
```

```
Kubectl rollout history deployment deployment/myapp-deployment
```

If something goes wrong, we can rollback to  
previous version:

```
Kubectl rollout undo deployment/myapp-deployment
```

# How to update a Deployment

- Update the container image to **adfonte/hello-server:newversion**.
  - First create the new image using the docker tag command based on adfonte/hello-server:latest

```
[afs-MacBook-Air:kubernetes alexandrefonte$ docker tag adfonte/hello-server:latest adfonte/hello-server:version
[afs-MacBook-Air:kubernetes alexandrefonte$ docker image ls -a
REPOSITORY          TAG      IMAGE ID   CREATED    SIZE
adfonte/hello-server  latest   d09ee31d0506  29 hours ago  544MB
adfonte/hello-server  newversion  d09ee31d0506  29 hours ago  544MB
adfonte/hello-server  version   d09ee31d0506  29 hours ago  544MB
product-server        latest   93310c05cbd6  6 days ago   491MB
message-server        latest   60ce7822e7c8  6 days ago   544MB
httpd                 latest   c30a46771695  9 days ago   144MB
kicbase/stable        v0.0.30  1312ccd2422d  2 months ago  1.14GB
```

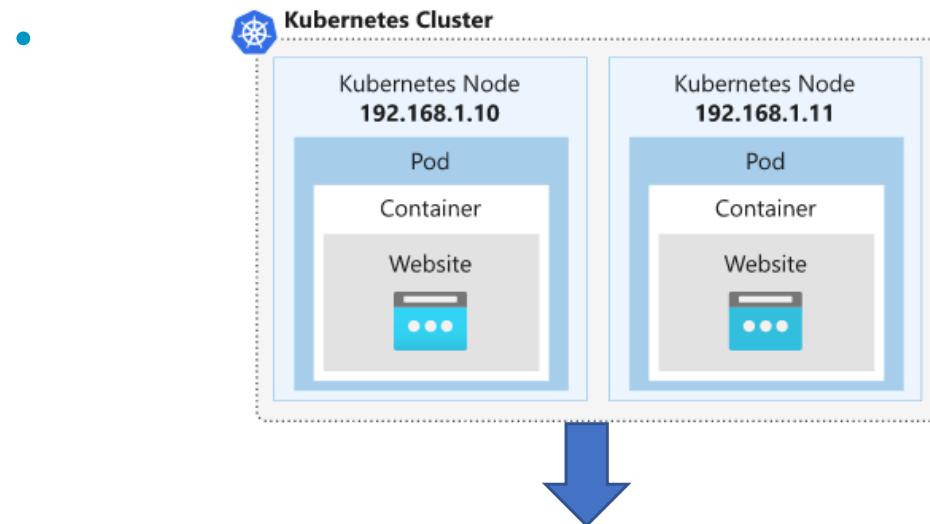
-Update the deployment

Result:

```
afs-MacBook-Air:kubernetes alexandrefonte$ kubectl set image deployment/hello-world-rest-api-deployment hello-world-rest-api=hello-world:newversion
deployment.apps/hello-world-rest-api-deployment image updated
```

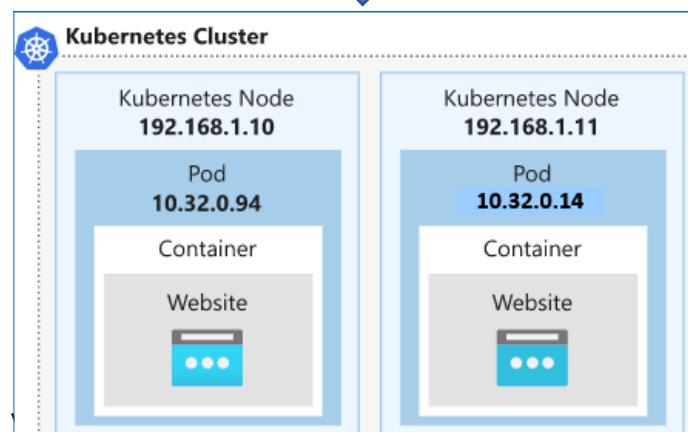
# Kubernetes networking

- Os nós têm endereços IP da rede onde estão ligados. Este endereço IP permite-nos o acesso ao nó via rest ou SSH.
- A cada pod é atribuído um endereço IP dinâmico (interno) de uma rede virtual separada da rede dos nós criada para ligar os pods



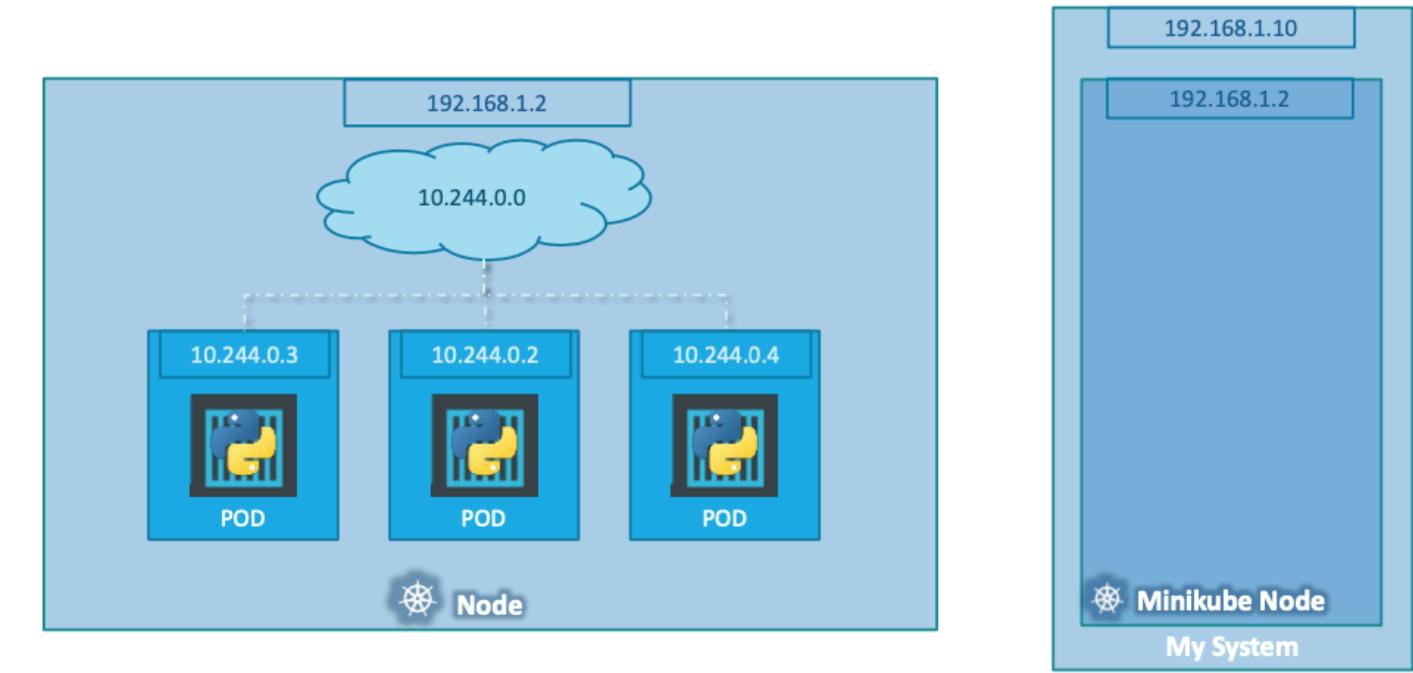
Although each Pod has a unique IP address, those IPs are not exposed outside the cluster without a Service.

Services allow your applications to receive traffic. Services can be exposed in different ways by specifying a type in the ServiceSpec



# Kubernetes networking with Minikube

- Host: 192.168.1.10
- Minikube Kubernetes node: 192.168.1.2 (virtualization)
- Pods: 10.244.0.(2-4)



# Service Types

- Although each Pod has a unique IP address, those IPs are not exposed outside the cluster without a Service. Services allow your applications to receive traffic. Services can be exposed in different ways by specifying a type in the ServiceSpec:
  - **ClusterIP (default)** - Exposes the Service on an internal IP in the cluster. This type makes the Service only reachable from within the cluster.
  - **NodePort** - Exposes the Service on the same port of each selected Node in the cluster using NAT. Makes a Service accessible from outside the cluster using <NodeIP>:<NodePort>. Superset of ClusterIP.
  - **LoadBalancer** - Creates an external load balancer in the current cloud (if supported) and assigns a fixed, external IP to the Service. Superset of NodePort.
  - **ExternalName** - Maps the Service to the contents of the externalName field (e.g. foo.bar.example.com), by returning a CNAME record with its value. No proxying of any kind is set up. This type requires v1.7 or higher of kube-dns, or CoreDNS version 0.0.8 or higher.

# Service Types - Cluster IP

- Service is like a virtual service inside the node ...

**-ClusterIP:** Exposes a service which is only accessible from within the cluster.

```
service-definition.yml
apiVersion: v1
kind: Service
metadata:
  name: back-end
spec:
  type: ClusterIP
  ports:
    - targetPort: 80
      port: 80
  selector:
```

```
pod-definition.yml
> kubectl create -f service-definition.yml
service "back-end" created

> kubectl get services
NAME        TYPE      CLUSTER-IP      EXTERNAL-IP   PORT(S)      AGE
kubernetes  ClusterIP  10.96.0.1     <none>       443/TCP     16d
back-end    ClusterIP  10.106.127.123  <none>       80/TCP      2m

          app: myapp
          type: back-end
spec:
  containers:
    - name: nginx-container
      image: nginx
```

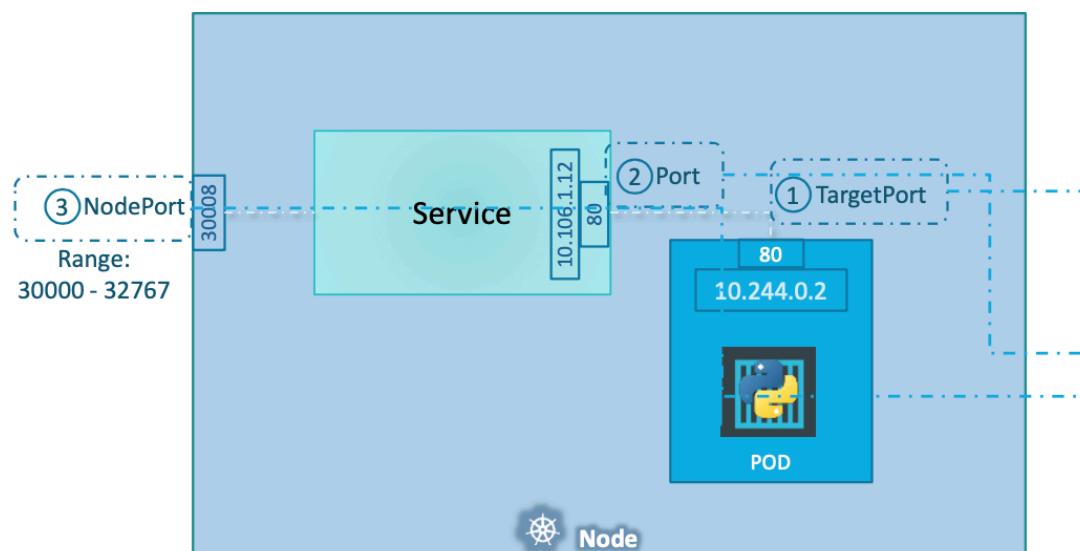
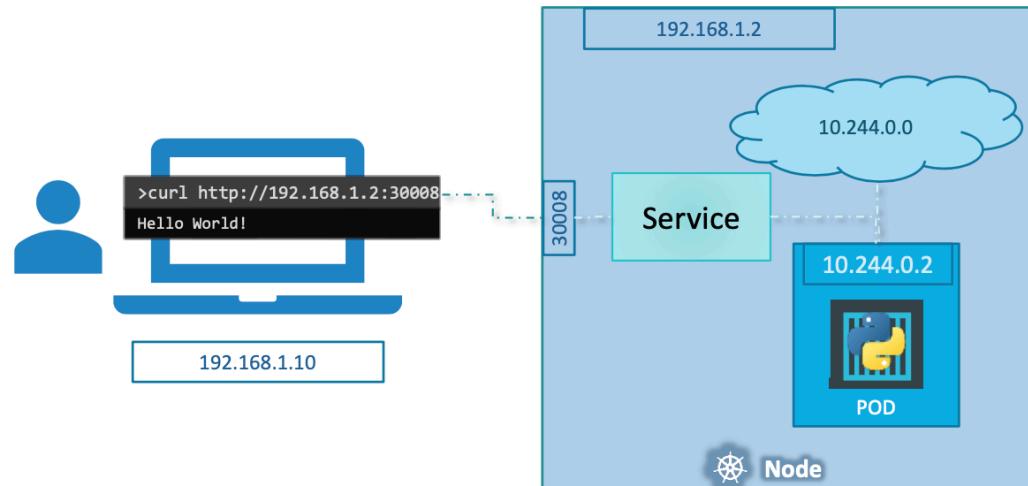
# NodePort vs LoadBalancer

- A **NodePort** service is the most basic way to get external traffic directly to your service. NodePort, as the name implies, opens a specific port, and any traffic that is sent to this port is forwarded to the service.
- A **LoadBalancer** service is the standard way to expose a service to the internet. With this method, each service gets its own IP address.
- Using minikube tunnel
  - Services of type LoadBalancer can be exposed via the minikube tunnel command. It must be run in a separate terminal window to keep the LoadBalancer running. Ctrl-C in the terminal can be used to terminate the process at which time the network routes will be cleaned up.

# Service Types - NodePort (External Communication)

- Service is like a virtual service inside the node ...

**-NodePort:** makes a internal pod accessible on a Port of the Node



There are

3 ports involved.

- Pod port (targetPort) where the container is running.
- Port of service itself
- Port of the node (range 30000-32767)

```
service-definition.yml
```

```
apiVersion: v1
kind: Service
metadata:
  name: myapp-service
spec:
  type: NodePort
  ports:
    - targetPort: 80
      *port: 80
      nodePort: 30008
```

# Service Types - NodePort (External Communication)

- Exemplo

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: my-nginx-deploy
  labels:
    app: nginx
spec:
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: test-nginx
          image: nginx:alpine
          ports:
            - containerPort: 80
replicas: 2
selector:
  matchLabels:
    app: nginx
```

```
kubectl apply -f nginx-deploy.yml
```

```
apiVersion: v1
kind: Service
metadata:
  name: my-service
spec:
  type: NodePort
  selector:
    app: nginx
  ports:
    - port: 80
      targetPort: 80
      nodePort: 30007
```

```
kubectl apply -f nodeport.yml
```

# Service Types - NodePort (External Communication)

- Create a service named **my-service** of NodePort type
- Create a port-forward from host computer port 7080 to 8080

```
afs-MacBook-Air:kubernetes alexandrefonte$ kubectl apply -f nodeport.yaml
The Service "my-service" is invalid: spec.ports[0].nodePort: Invalid value: 7080: provided port is
not in the valid range. The range of valid ports is 30000-32767
afs-MacBook-Air:kubernetes alexandrefonte$ kubectl apply -f nodeport.yaml
service/my-service created
afs-MacBook-Air:kubernetes alexandrefonte$ kubectl get all
NAME                                         READY   STATUS    RESTARTS   AGE
pod/hello-world-rest-api-deployment-7d645f687f-bsq67   1/1     Running   0          3m23s
pod/hello-world-rest-api-deployment-7d645f687f-dwkgw   1/1     Running   0          3m23s
pod/hello-world-rest-api-deployment-7d645f687f-twsss   1/1     Running   0          3m23s

NAME           TYPE        CLUSTER-IP      EXTERNAL-IP      PORT(S)          AGE
service/kubernetes   ClusterIP   10.96.0.1    <none>           443/TCP         8d
service/my-service   NodePort    10.110.2.30  <none>           8080:30000/TCP  5s

NAME                           READY   UP-TO-DATE   AVAILABLE   AGE
deployment.apps/hello-world-rest-api-deployment   3/3     3           3           3m24s

NAME                           DESIRED   CURRENT   READY   AGE
replicaset.apps/hello-world-rest-api-deployment-7d645f687f   3        3        3       3m24s
```

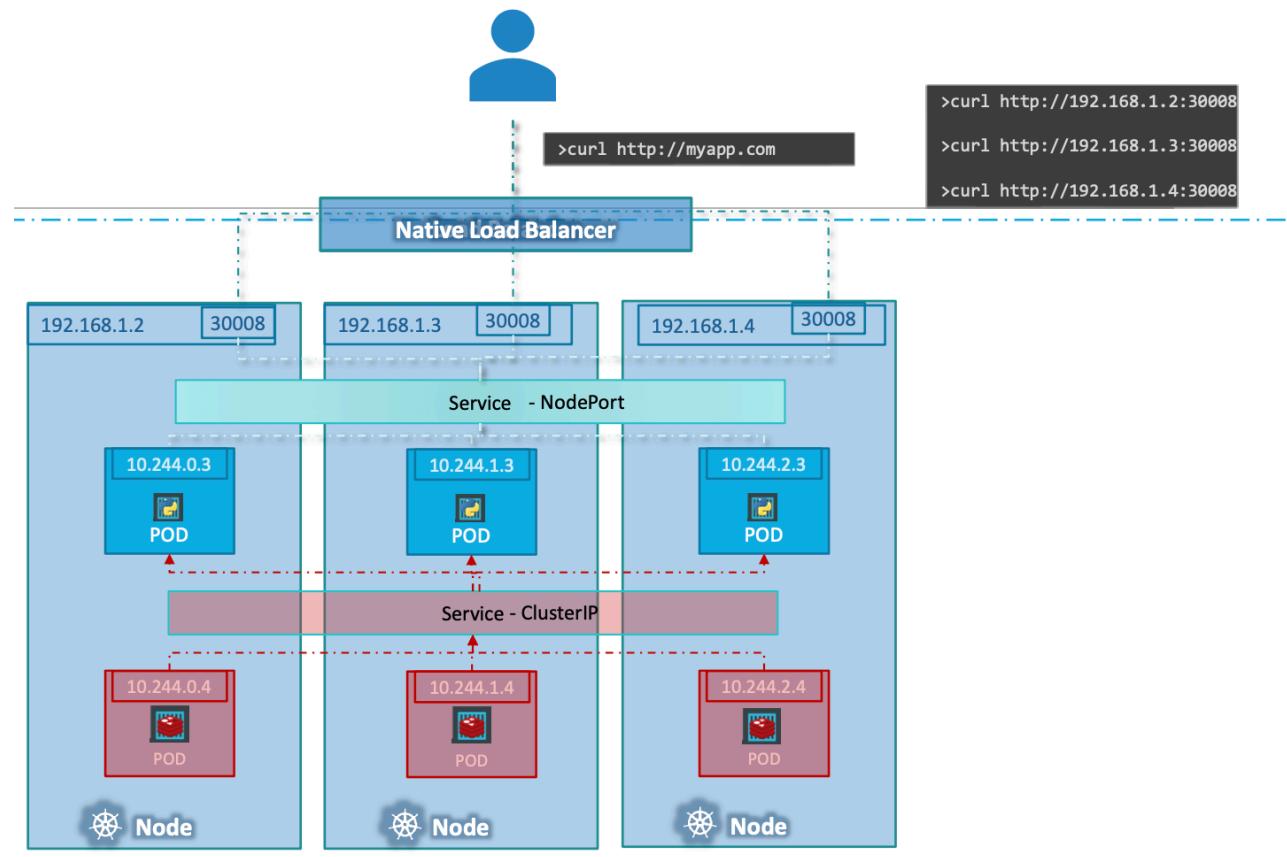
```
[afs-MacBook-Air:kubernetes alexandrefonte$ kubectl port-forward deployment/hello-world-rest-api-deployment 7080:8080
Forwarding from 127.0.0.1:7080 -> 8080
Forwarding from [::1]:7080 -> 8080
```

## nodeport.yaml

```
apiVersion: v1
kind: Service
metadata:
  name: my-service
spec:
  type: NodePort
  selector:
    app: hello-world-rest-api
  ports:
    - port: 8080
      targetPort: 8080
      nodePort: 30000
```

# Service Types - LoadBalancer (External Communication)

- Service is like a virtual service inside the node ...  
**-LoadBalancer** service type is built on top of NodePort service types by provisioning and configuring external load balancers. It exposes services that are running in the cluster by forwarding layer 4 traffic to worker nodes.



# II) Service Types - LoadBalancer (External Communication)

- Service is like a virtual service inside the node ...  
**-LoadBalancer (Cont.)**

```
service-definition.yml
```

```
apiVersion: v1
kind: Service
metadata:
  name: front-end
spec:
  type: NodeBalancer
  ports:
    - targetPort: 80
      port: 80
  selector:
    app: myapp
    type: front-end
```

```
> kubectl create -f service-definition.yml
service "front-end" created
```

```
> kubectl get services
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	16d
front-end	LoadBalancer	10.106.127.123	<Pending>	80/TCP	2m

# Service Types - LoadBalancer (External Communication)

- Create a service named **my-service-lb** of LoadBalancer type  
-nodePort: 30001

```
[afs-MacBook-Air:kubernetes alexandrefonte$ kubectl get service
NAME      TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
kubernetes  ClusterIP  10.96.0.1    <none>        443/TCP     8d
my-service  NodePort   10.110.2.30  <none>        8080:30000/TCP 22m
[afs-MacBook-Air:kubernetes alexandrefonte$ kubectl apply -f loadbalancer.yaml
service/my-service-lb created
[afs-MacBook-Air:kubernetes alexandrefonte$ kubectl get service
NAME      TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
kubernetes  ClusterIP  10.96.0.1    <none>        443/TCP     8d
my-service  NodePort   10.110.2.30  <none>        8080:30000/TCP 22m
my-service-lb  LoadBalancer  10.102.128.100  <pending>    8080:30001/TCP 2s
```

loadbalancer.yaml

```
apiVersion: v1
kind: Service
metadata:
  name: my-service-lb
spec:
  type: LoadBalancer
  selector:
    app: hello-world-rest-api
  ports:
# Port accessible inside cluster
    - port: 8080
# Port to forward to inside the pod
    targetPort: 8080
# Port accessible outside cluster
    nodePort: 30001
```

```
[afs-MacBook-Air:kubernetes alexandrefonte$ minikube tunnel
✓ Tunnel successfully started
✗ NOTE: Please do not close this terminal as this process must stay alive for the tunnel to be accessible ...
🏃 Starting tunnel for service my-service-lb.
```

```
[afs-MacBook-Air:kubernetes alexandrefonte$ kubectl get service
NAME      TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
kubernetes  ClusterIP  10.96.0.1    <none>        443/TCP     8d
my-service  NodePort   10.110.2.30  <none>        8080:30000/TCP 24m
my-service-lb  LoadBalancer  10.102.128.100  127.0.0.1    8080:30001/TCP 103s
afs-MacBook-Air:kubernetes alexandrefonte$
```

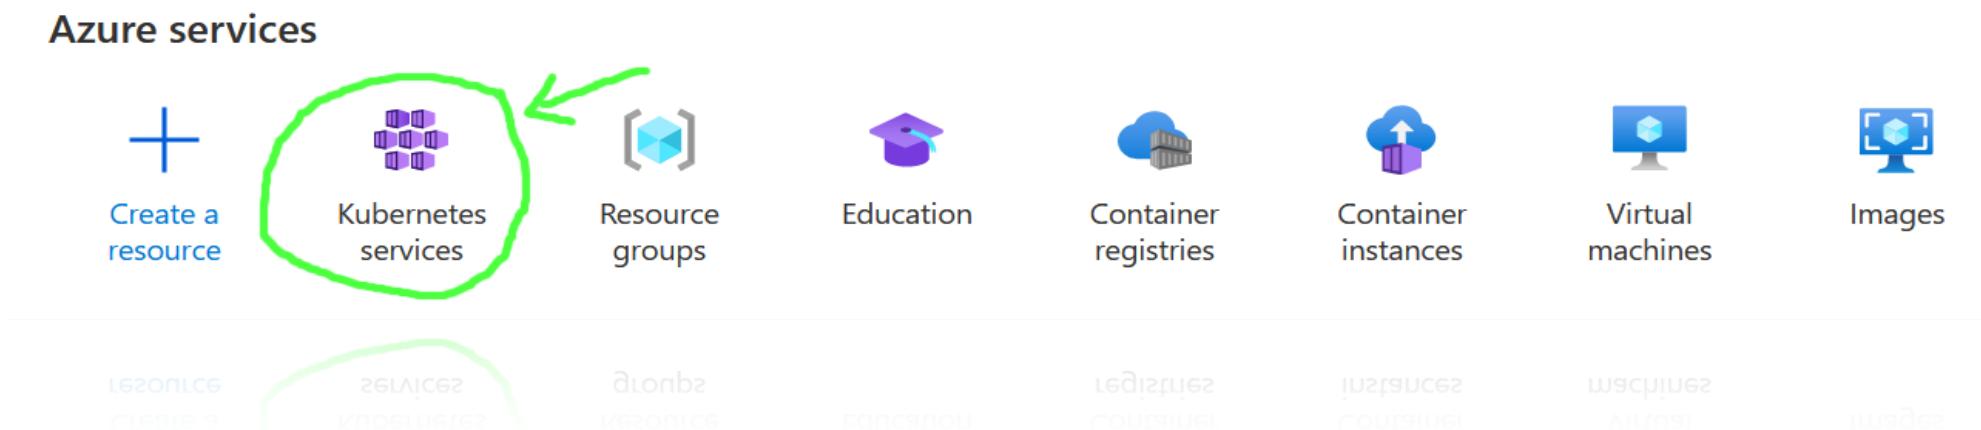
In a different terminal window

In case of Minikube running as a VM, appears a different IP

In a browser enter:  
127.0.0.1:8080/messages/students/aid

# Azure Kubernetes Service (AKS)

- O serviço AKS providencia um ambiente kubernetes integrado na cloud Azure
- Desta forma evita-se todo o trabalho adicional associado à implementação de kubernetes em ambientes físicos
- O control plane (master node) é gratuito, só se paga as VM onde vão ficar os outros nós (worker nodes), storage e networking



# Criação de um cluster K8s no AKS

## Create Kubernetes cluster

Basics   Node pools   Authentication   Networking   Integrations   Tags   Review + create

Azure Kubernetes Service (AKS) manages your hosted Kubernetes environment, making it quick and easy to deploy and manage containerized applications without container orchestration expertise. It also eliminates the burden of ongoing operations and maintenance by provisioning, upgrading, and scaling resources on demand, without taking your applications offline. [Learn more about Azure Kubernetes Service](#)

**Project details**

Select a subscription to manage deployed resources and costs. Use resource groups like folders to organize and manage all your resources.

Subscription \* ⓘ    ▼

Resource group \* ⓘ    ▼  
[Create new](#)

**Cluster details**

Kubernetes cluster name \* ⓘ    ✓

Region \* ⓘ    ▼

Availability zones    ▼

Kubernetes version \* ⓘ    ▼

**Primary node pool**

The number and size of nodes in the primary node pool in your cluster. For production workloads, at least 3 nodes are recommended for resiliency. For development or test workloads, only one node is required. If you would like to add additional node pools or to see additional configuration options for this node pool, go to the 'Node pools' tab above. You will be able to add additional node pools after creating your cluster. [Learn more about node pools in Azure Kubernetes Service](#)

Node size \* ⓘ    ▼  
[Change size](#)

Node count \* ⓘ    ▼

# Acesso ao cluster criado por AKS

- A interação com o cluster criado no AKS pode ser feita da forma habitual - usando o comando CLI kubectl
- O kubectl já está integrado no Azure Cloud Shell
- Mas também pode ser instalado localmente

Connect to myKubernetesCluster X

Connect to your cluster using command line tooling to interact directly with cluster using kubectl, the command line tool for Kubernetes. Kubectl is available within the Azure Cloud Shell by default and can also be installed locally. [Learn more ↗](#)

1. [Open Cloud Shell ↗](#) or the Azure CLI  
2. Run the following commands

```
az account set --subscription e38bf7d7-9e4d-4832-810a-88a6b884ab64
```

```
az aks get-credentials --resource-group testes --name myKubernetesCluster
```

**Sample commands**  
Once you have run the command above to connect to the cluster, you can run any kubectl commands. Here are a few examples of useful commands you can try.

```
# List all deployments in all namespaces
kubectl get deployments --all-namespaces=true
```

```
# List all deployments in a specific namespace
# Format :kubectl get deployments --namespace <namespace-name>
kubectl get deployments --namespace kube-system
```

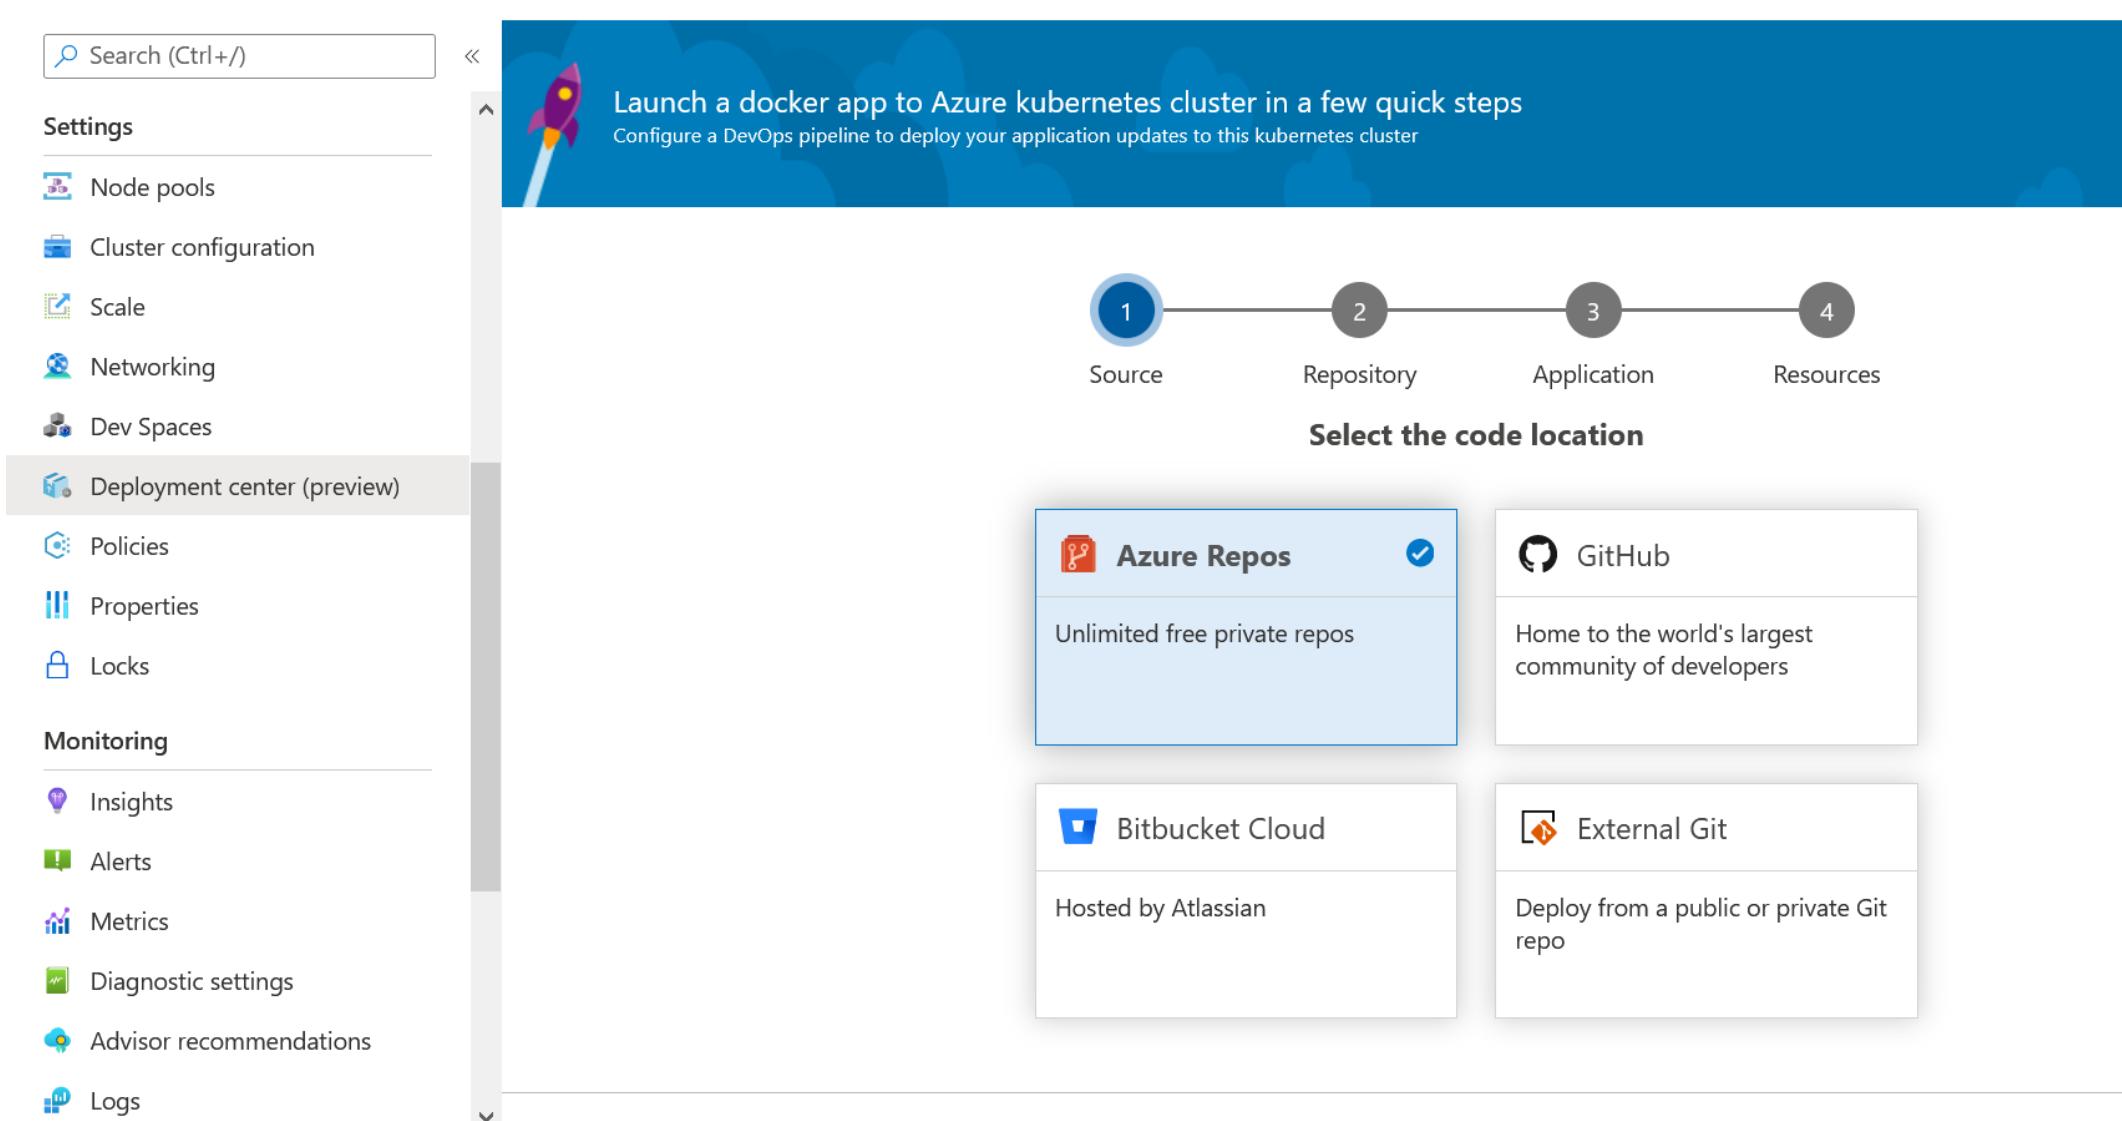
```
# List details about a specific deployment
# Format :kubectl describe deployment <deployment-name> --namespace
<namespace-name>
kubectl describe deployment my-dep --namespace kube-system
```

```
# List pods using a specific label
# Format :kubectl get pods -l <label-key>=<label-value> --all-namespaces=true
kubectl get pods -l app=nginx --all-namespaces=true
```

```
# Get logs for all pods with a specific label
# Format :kubectl logs -l <label-key>=<label-value>
kubectl logs -l app=nginx --namespace kube-system
```

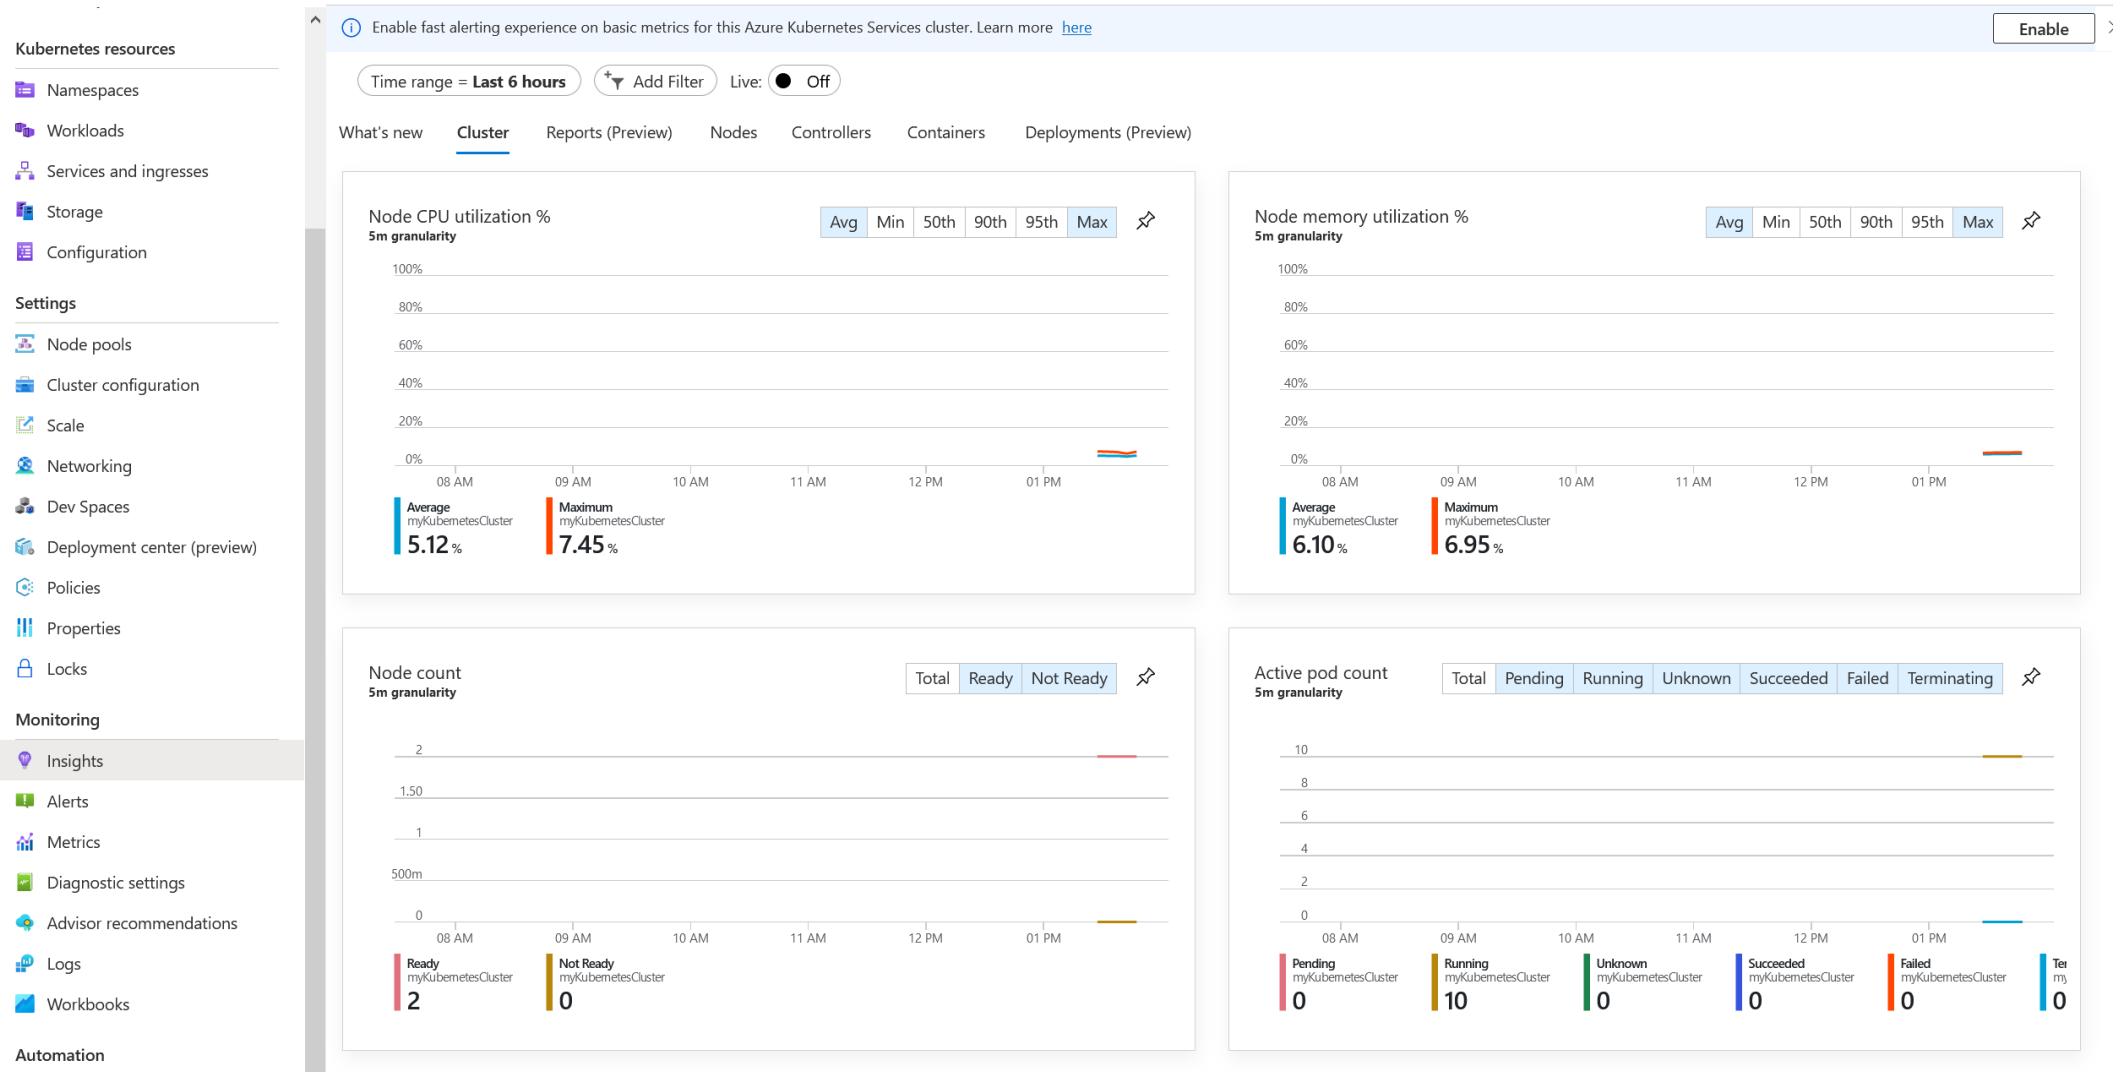
# Também pode ser usado o portal Azure...

- Nesta imagem pode ser visto o deployment center

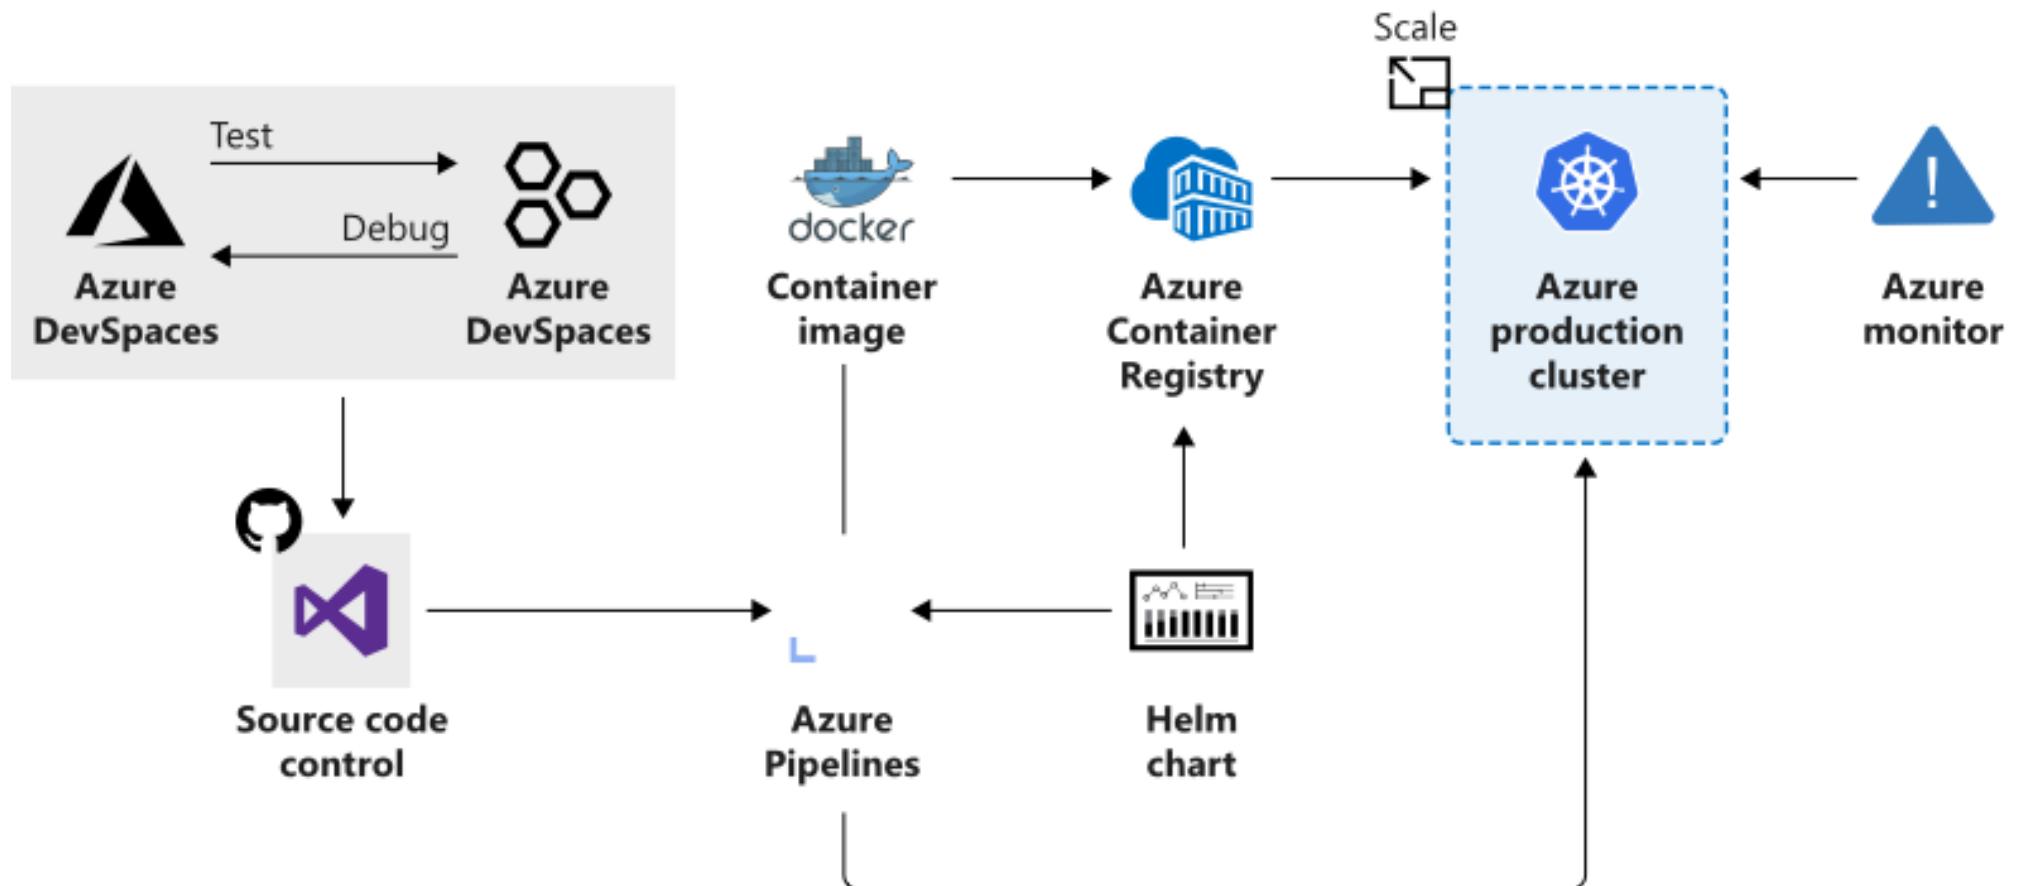


The screenshot shows the Azure Deployment Center (preview) interface. On the left, there's a navigation sidebar with various options like Settings, Node pools, Cluster configuration, Scale, Networking, Dev Spaces, Deployment center (preview), Policies, Properties, Locks, Monitoring, Insights, Alerts, Metrics, Diagnostic settings, Advisor recommendations, and Logs. The 'Deployment center (preview)' option is highlighted. The main area has a title 'Launch a docker app to Azure kubernetes cluster in a few quick steps' with a subtitle 'Configure a DevOps pipeline to deploy your application updates to this kubernetes cluster'. Below this is a horizontal step bar with four circles labeled 1 through 4, corresponding to 'Source', 'Repository', 'Application', and 'Resources'. A call-to-action button 'Select the code location' is centered below the step bar. There are four options listed: 'Azure Repos' (selected, indicated by a checked checkbox icon), 'GitHub', 'Bitbucket Cloud', and 'External Git'. Each option has a brief description: 'Unlimited free private repos' for Azure Repos, 'Home to the world's largest community of developers' for GitHub, 'Hosted by Atlassian' for Bitbucket Cloud, and 'Deploy from a public or private Git repo' for External Git.

# Monitorização do cluster kubernetes



# Desenvolvimento e implementação de workloads



# Questões

